



Speedrun through Splicing Sockets with SOCKMAP

Jakub Sitnicki
Systems Engineer
Cloudflare



NDC { TechTown }

12 Sep 2024 | Kongsberg, Norway



Walk-through of ~~Speedrun through~~ Splicing Sockets with SOCKMAP

Jakub Sitnicki
Systems Engineer
Cloudflare



**DIRECTOR'S
CUT**

NDC { TechTown }

12 Sep 2024 | Kongsberg, Norway

\$ whoami

Linux / OS Team @ Cloudflare

🌽 roll out fresh kernels

🐛 squash bugs





🧑🔧 troubleshoot stuff

✨ prototype features







\$ whoami

Linux / OS Team @ Cloudflare

-  roll out fresh kernels
-  squash bugs
-  troubleshoot stuff
-  prototype features

SOCKMAP co-maintainer @ Linux upstream

-  small-time (= feature) maintainer
-  fix bugs
-  review patches
-  answer questions



About this talk

Good to know:

- ❑ network programming (`socket`, `connect`, `sendmsg`, `recvmsg`)
- ❑ basics of eBPF (what are BPF maps, programs, hooks, `bpftool`)
- ❑ building blocks of containers (cgroups, namespaces)

Goals:

- ❑ know that SOCKMAP exists
- ❑ have idea how / when / what for you can use it
- ❑ feel ready to dive deeper

Agenda

- 1 What can SOCKMAP do?
- 2 What is SOCKMAP?
- 3 Evolution of SOCKMAP
- 4 How to set up SOCKMAP?
- 5 How to get sockets into a SOCKMAP?
- 6 Supported socket splicing setups
- 7 Real-life use cases



**REST AREA
1 MILE**

Agenda

- 1 What can SOCKMAP do?
- 2 What is SOCKMAP?
- 3 Evolution of SOCKMAP
- 4 How to set up SOCKMAP?
- 5 How to get sockets into a SOCKMAP?
- 6 Supported socket splicing setups
- 7 Real-life use cases

NEW STUFF

- * more benchmarks
- * internal design
- * how to trace it

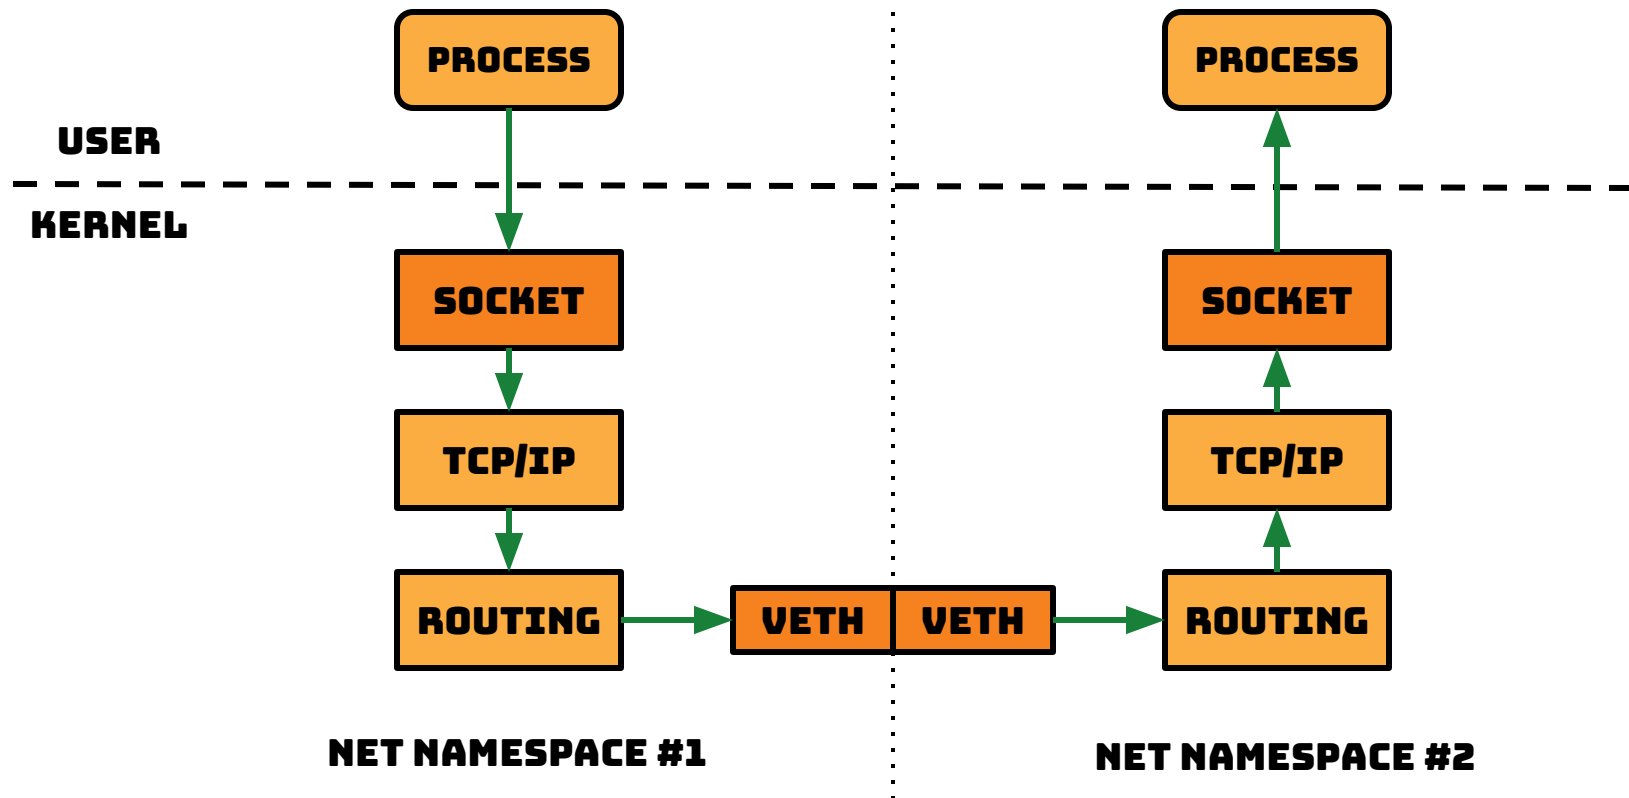


REST AREA
1 MILE

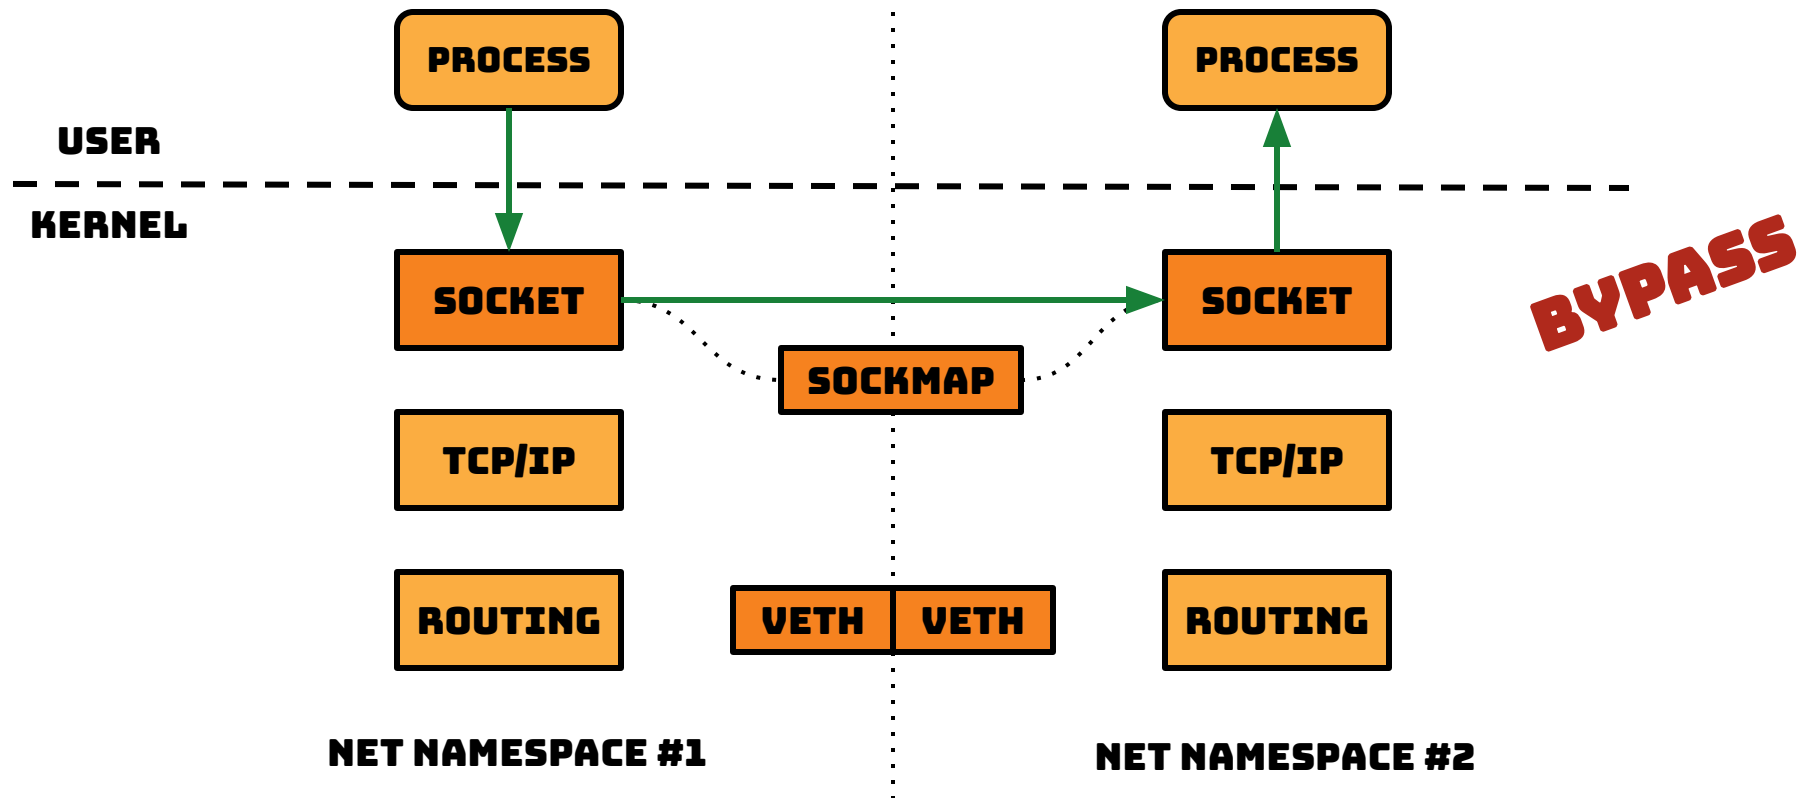
What can SOCKMAP do for you?



What can SOCKMAP do for... container networking



What can SOCKMAP do for... container networking



Let's set it up!


Create two network namespaces

```
# ip netns add A
```

```
# ip netns add B
```

NETNS A

NETNS B



Let's set it up!


Create two network namespaces

```
# ip netns add A
```

```
# ip netns add B
```

Link network namespaces with a veth pair

```
# ip -n A link add name veth1 type veth peer name veth2 netns B
```



Let's set it up!

Create two network namespaces

```
# ip netns add A
```

```
# ip netns add B
```


Link network namespaces with a veth pair

```
# ip -n A link add name veth1 type veth peer name veth2 netns B
```

Bring up the links inside network namespaces

```
# ip -n A link set dev veth1 up
```

```
# ip -n B link set dev veth2 up
```



Let's set it up!

Create two network namespaces

```
# ip netns add A
```

```
# ip netns add B
```

Link network namespaces with a veth pair

```
# ip -n A link add name veth1 type veth peer name veth2 netns B
```

Bring up the links inside network namespaces

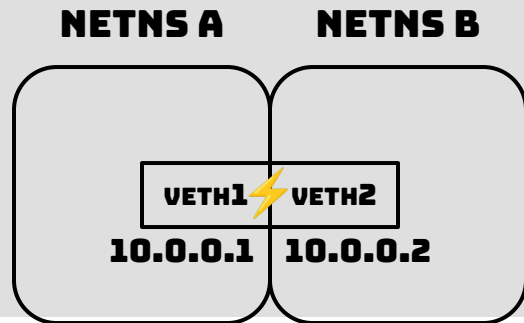
```
# ip -n A link set dev veth1 up
```

```
# ip -n B link set dev veth2 up
```

Assign addresses to links inside network namespaces

```
# ip -n A addr add 10.0.0.1/24 dev veth1
```

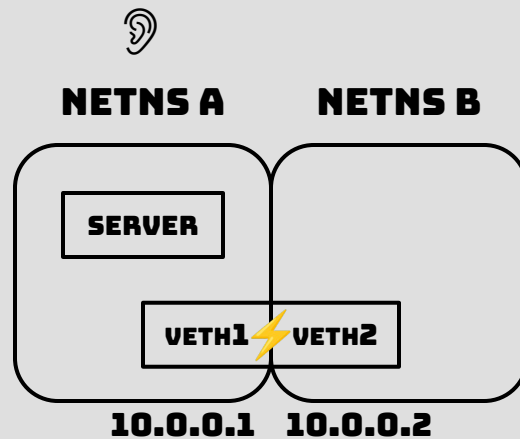
```
# ip -n B addr add 10.0.0.2/24 dev veth2
```



Measure latency, no SOCKMAP first

Run TCP server in netns A

```
# ip netns exec A \  
  sockperf server -i 10.0.0.1 --tcp --daemonize
```



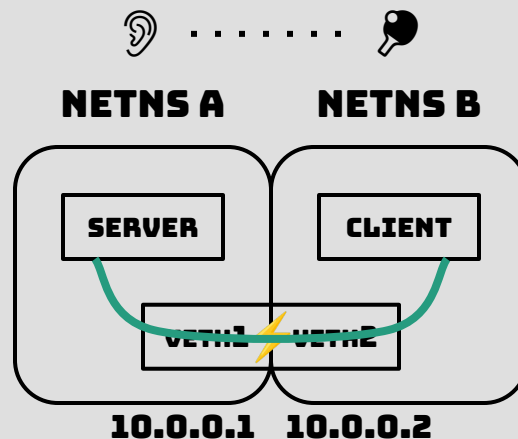
Measure latency, no SOCKMAP first

Run TCP server in netns A

```
# ip netns exec A \  
  sockperf server -i 10.0.0.1 --tcp --daemonize
```

Run TCP client in netns B

```
# ip netns exec B \  
  sockperf ping-pong -i 10.0.0.1 --tcp --time 30
```



Measure latency, no SOCKMAP first

Run TCP server in netns A

```
# ip netns exec A \  
  sockperf server -i 10.0.0.1 --tcp --daemonize
```

Run TCP client in netns B

```
# ip netns exec B \  
  sockperf ping-pong -i 10.0.0.1 --tcp --time 30
```

...

```
sockperf: [Total Run] RunTime=30.000 sec; Warm up time=400 msec; SentMessages=2599753;  
ReceivedMessages=2599752
```

...

```
sockperf: ==> avg-latency=5.748 (std-dev=2.010, mean-ad=0.322, median-ad=0.220,  
sigr=0.239, cv=0.350, std-error=0.001, 99.0% ci=[5.745, 5.751])  
sockperf: # dropped messages = 0; # duplicated messages = 0; # out-of-order messages = 0  
sockperf: Summary: Latency is 5.748 usec
```

$5.8 \pm 2.0 \mu\text{sec}$

Set up SOCKMAP bypass

Load BPF programs and create BPF maps

```
# bpftool prog loadall \  
    redir_bypass.bpf.o /sys/fs/bpf pinmaps /sys/fs/bpf
```


**BPF
PROGS**

SK_MSG

SK_OPS

BPF MAPS

SOCKMAP





Set up SOCKMAP bypass

Load BPF programs and create BPF maps

```
# bpftool prog loadall \  
    redir_bypass.bpf.o /sys/fs/bpf pinmaps /sys/fs/bpf
```

Attach BPF program to BPF map

```
# bpftool prog attach \  
    pinned /sys/fs/bpf/sk_msg_prog sk_msg_verdict \  
    pinned /sys/fs/bpf/sock_map
```



Set up SOCKMAP bypass

Load BPF programs and create BPF maps


```
# bpftool prog loadall \  
    redir_bypass.bpf.o /sys/fs/bpf pinmaps /sys/fs/bpf
```

Attach BPF program to BPF map

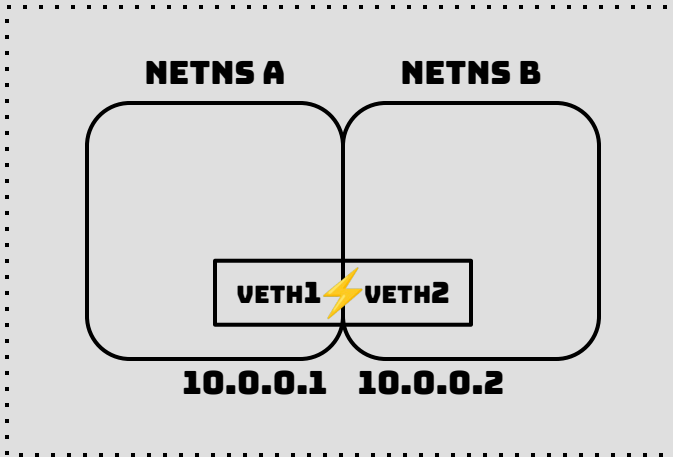
```
# bpftool prog attach \  
    pinned /sys/fs/bpf/sk_msg_prog sk_msg_verdict \  
    pinned /sys/fs/bpf/sock_map
```

Create a test cgroup

```
# mkdir /sys/fs/cgroup/test.slice
```



TEST.SLICE CGROUP



Set up SOCKMAP bypass

Load BPF programs and create BPF maps

```
# bpftool prog loadall redir_bypass.bpf.o /sys/fs/bpf \  
    pinmaps /sys/fs/bpf
```

Attach BPF program to BPF map

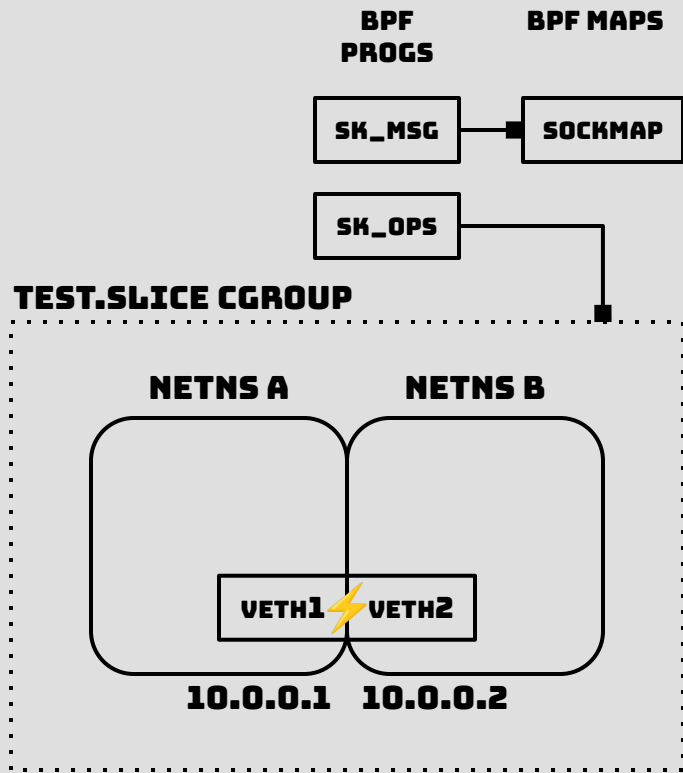
```
# bpftool prog attach \  
    pinned /sys/fs/bpf/sk_msg_prog sk_msg_verdict \  
    pinned /sys/fs/bpf/sock_map
```

Create a test cgroup

```
# mkdir /sys/fs/cgroup/test.slice
```

Attach BPF program to cgroup

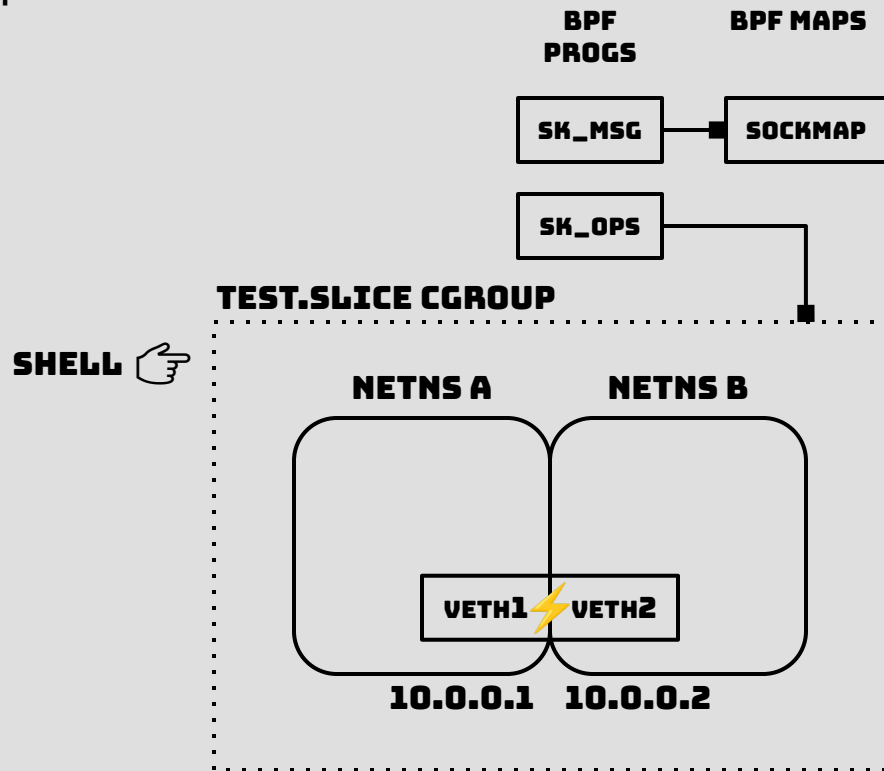
```
# bpftool cgroup attach \  
    /sys/fs/cgroup/test.slice \  
    cgroup_sock_ops pinned /sys/fs/bpf/sockops_prog
```



Repeat test with SOCKMAP bypass

Spawn client and server inside the test cgroup

```
# echo $$ > /sys/fs/cgroup/test.slice/cgroup.procs
```



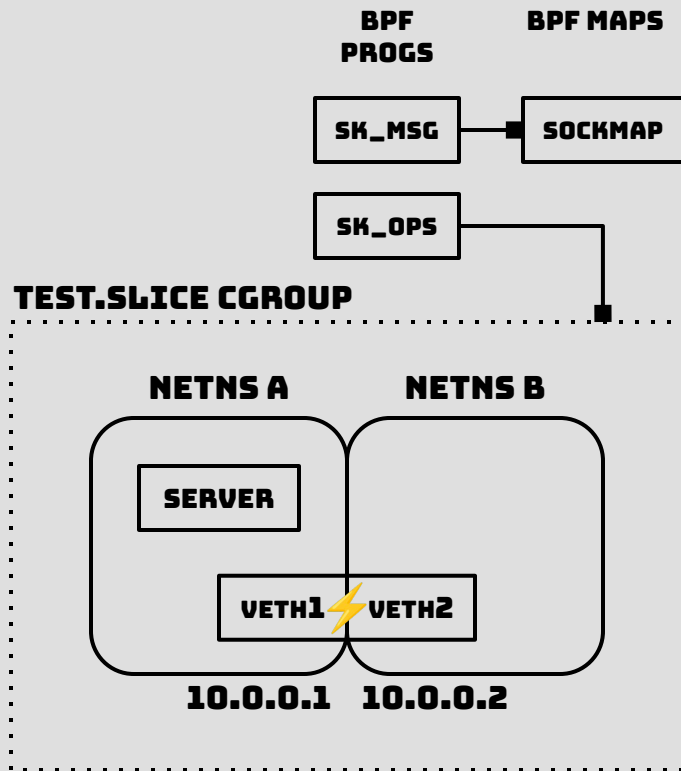
Repeat test with SOCKMAP bypass

Spawn client and server inside the test cgroup

```
# echo $$ > /sys/fs/cgroup/test.slice/cgroup.procs
```

Run TCP server in netns A

```
# ip netns exec A \  
  sockperf server -i 10.0.0.1 --tcp --daemonize
```



Repeat test with SOCKMAP bypass

Spawn client and server inside the test cgroup

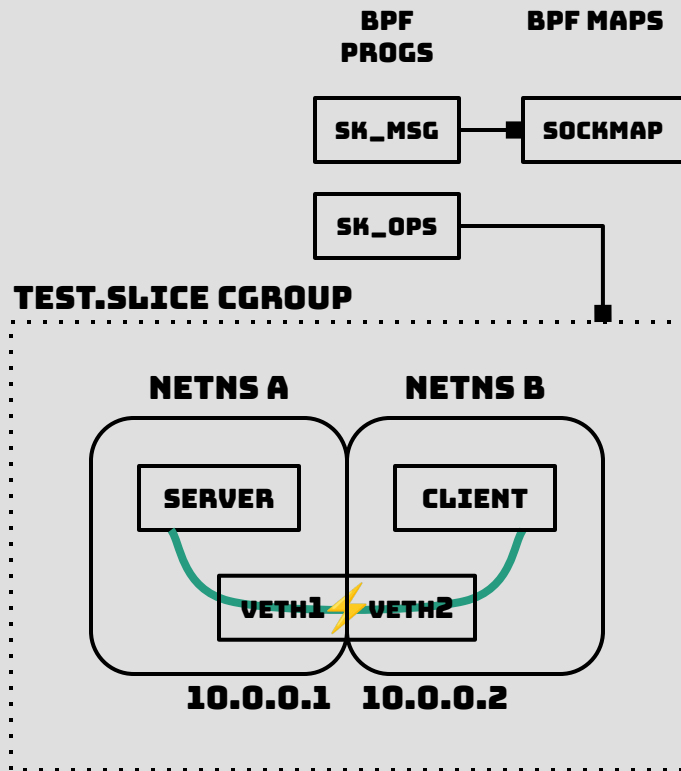
```
# echo $$ > /sys/fs/cgroup/test.slice/cgroup.procs
```

Run TCP server in netns A

```
# ip netns exec A \  
  sockperf server -i 10.0.0.1 --tcp --daemonize
```

Run TCP client

```
# ip netns exec B \  
  sockperf ping-pong -i 10.0.0.1 --tcp --time 30
```



Repeat test with SOCKMAP bypass

Spawn client and server inside the test cgroup

```
# echo $$ > /sys/fs/cgroup/test.slice/cgroup.procs
```

Run TCP server in netns A

```
# ip netns exec A \  
    sockperf server -i 10.0.0.1 --tcp --daemonize
```

Run TCP client

```
# ip netns exec B \  
    sockperf ping-pong -i 10.0.0.1 --tcp --time 30
```

```
sockperf: [Total Run] RunTime=30.000 sec; Warm up time=400 msec; SentMessages=3189584;  
ReceivedMessages=3189583
```

...

```
sockperf: ==> avg-latency=4.686 (std-dev=2.862, mean-ad=0.250, median-ad=0.216,  
siqr=0.173, cv=0.611, std-error=0.002, 99.0% ci=[4.682, 4.690])
```

```
sockperf: # dropped messages = 0; # duplicated messages = 0; # out-of-order messages = 0
```

```
sockperf: Summary: Latency is 4.686 usec
```

$4.7 \pm 2.9 \mu\text{sec}$

Without and with SOCKMAP bypass

before: $5.8 \pm 2.0 \mu\text{sec}$

↓ - 18.5%

after: $4.7 \pm 2.9 \mu\text{sec}$

Run the benchmark yourself:

<https://github.com/jsitnicki/sockmap-project/tree/main/examples/send-to-local>

What is SOCKMAP?




What is SOCKMAP?



Two things

What is SOCKMAP?

**Collection / container
for socket references
in Linux kernel**



What is SOCKMAP?


SOCKMAP
API



1. container for sockets


- BPF map (K/V store)
- holds weak refs to sockets

What is SOCKMAP?



API for enforcing policy
and
redirecting data between sockets

What is SOCKMAP?



EVOLUTION OF SOCKMAP



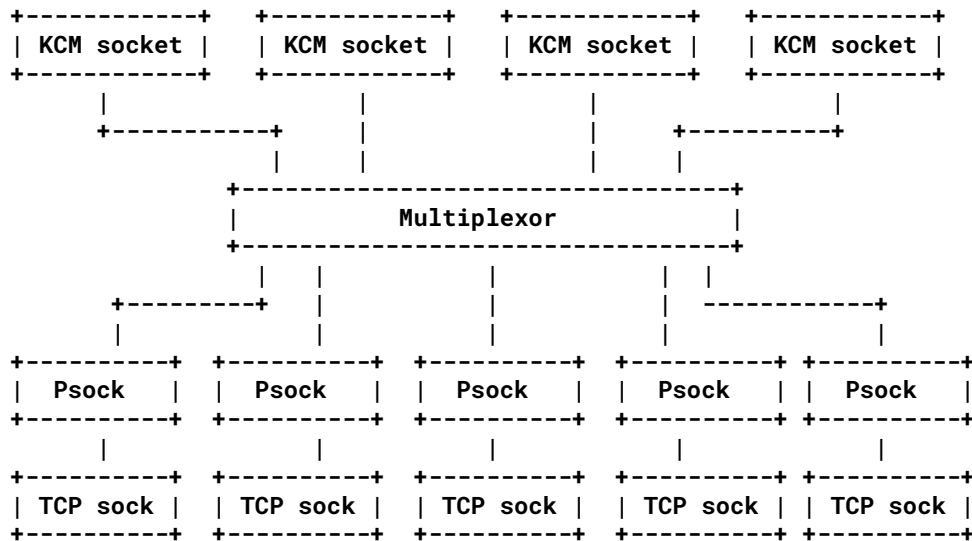
How did we get here?

Kernel Connection Multiplexor

Added infrastructure later reused by SOCKMAP - Psock and stream parser program.

2016

v4.6



How did we get here?



Kernel Connection Multiplexor

Added infrastructure later reused by SOCKMAP - Psock and stream parser program.

2016

v4.6

2017

v4.14

SOCKMAP initial version

Filtering and redirect on ingress to socket layer.

Counterpart of XDP DEVMAP.

```
author      John Fastabend <john.fastabend@gmail.com> 2017-08-15 22:32:47 -0700
committer   David S. Miller <davem@davemloft.net> 2017-08-16 11:27:53 -0700
commit      174a79ff9515f400b9a6115643dafd62a635b7e6 (patch)
tree        f48f1fc407adb9bce6fb0e5cddaabd7141acd071
parent      a6f6df69c48b86cd84f36c70593eb4968fceb34a (diff)
download    linux-174a79ff9515f400b9a6115643dafd62a635b7e6.tar.gz
```

bpf: sockmap with sk redirect support

Recently we added a new map type called dev map used to forward XDP packets between ports (6093ec2dc313). This patches introduces a similar notion for sockets.

A sockmap allows users to add participating sockets to a map. When sockets are added to the map enough context is stored with the map entry to use the entry with a new helper

```
bpf_sk_redirect_map(map, key, flags)
```

This helper (analogous to `bpf_redirect_map` in XDP) is given the map and an entry in the map. When called from a sockmap program, discussed below, the skb will be sent on the socket using `skb_send_sock()`.

commit 174a79ff9515 ("bpf: sockmap with sk redirect support")

How did we get here?



Kernel Connection Multiplexor

Added infrastructure later reused by SOCKMAP - Psock and stream parser program.

2016

v4.6

2017

v4.14

2018

v4.17

SK_MSG program

Filter and redirect at sendmsg() time (TCP only).

SOCKMAP initial version

Filter and redirect on ingress to socket layer (TCP only).

Counterpart of XDP DEVMAP.

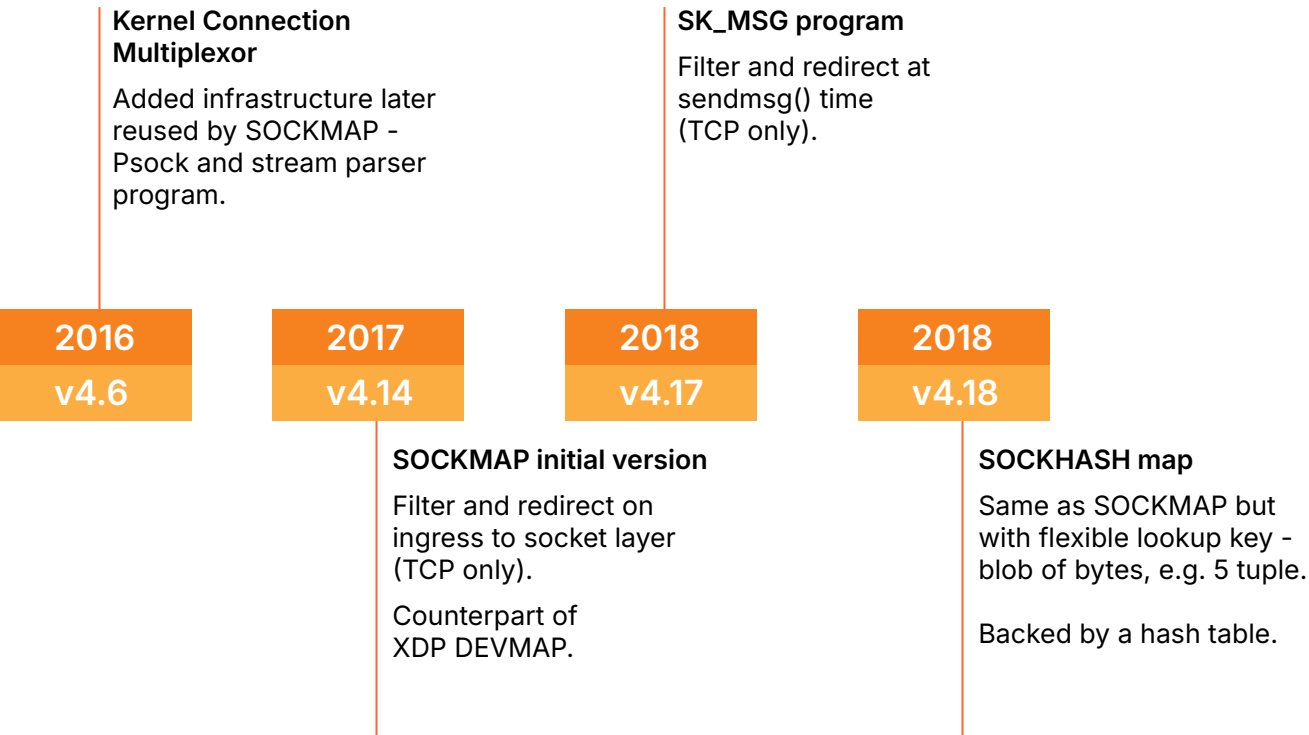
```
author      John Fastabend <john.fastabend@gmail.com> 2018-03-18 12:57:10 -0700
committer   Daniel Borkmann <daniel@iogearbox.net> 2018-03-19 21:14:38 +0100
commit      4f738adba30a7cfc006f605707e7aee847ffefa0 (patch)
tree        6603749a44356d3a44110c44f890a45b88d7e935
parent      8c05dbf04b2882c3c0bc43fe7668c720210877f3 (diff)
download    linux-4f738adba30a7cfc006f605707e7aee847ffefa0.tar.gz
```

bpf: create tcp_bpf_ulp allowing BPF to monitor socket TX/RX data

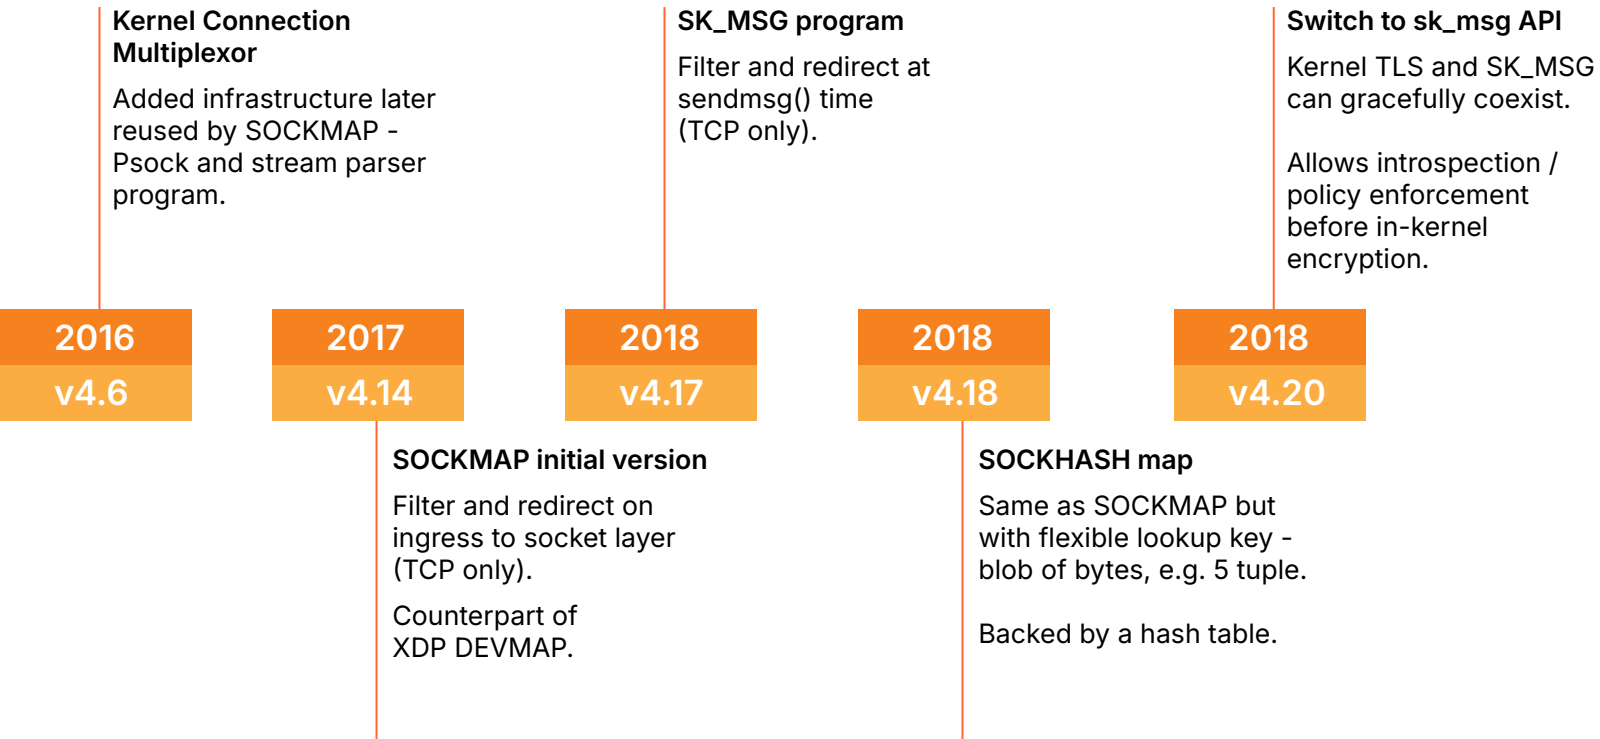
This implements a BPF ULP layer to allow policy enforcement and monitoring at the socket layer. In order to support this a new program type BPF_PROG_TYPE_SK_MSG is used to run the policy at the sendmsg/sendpage hook. To attach the policy to sockets a sockmap is used with a new program attach type BPF_SK_MSG_VERDICT.

commit 4f738adba30a ("bpf: create tcp_bpf_ulp allowing BPF to monitor socket TX/RX data")

How did we get here?



How did we get here?



commit 4f738adba30a ("bpf: create tcp_bpf_ulp allowing BPF to monitor socket TX/RX data")

How did we get here?

Store TCP or UDP sockets

SOCKMAP becomes a generic BPF map for sockets.

It can hold both connected and listening TCP sockets, and any bound UDP socket.

2020

v5.7

How did we get here?

Store TCP or UDP sockets

SOCKMAP becomes a generic BPF map for sockets.

It can hold both connected and listening TCP sockets, and any bound UDP socket.

2020

v5.7

2020

v5.10

BPF iterators support

Iterate over SOCKMAP from BPF context.

Allows copying socket references from one SOCKMAP to another.

How did we get here?

Store TCP or UDP sockets

SOCKMAP becomes a generic BPF map for sockets.

It can hold both connected and listening TCP sockets, and any bound UDP socket.

2020

v5.7

2020

v5.10

Redirect for UDP

Use new SK_SKB_VERDICT BPF program to redirect packets between UDP sockets.

2021

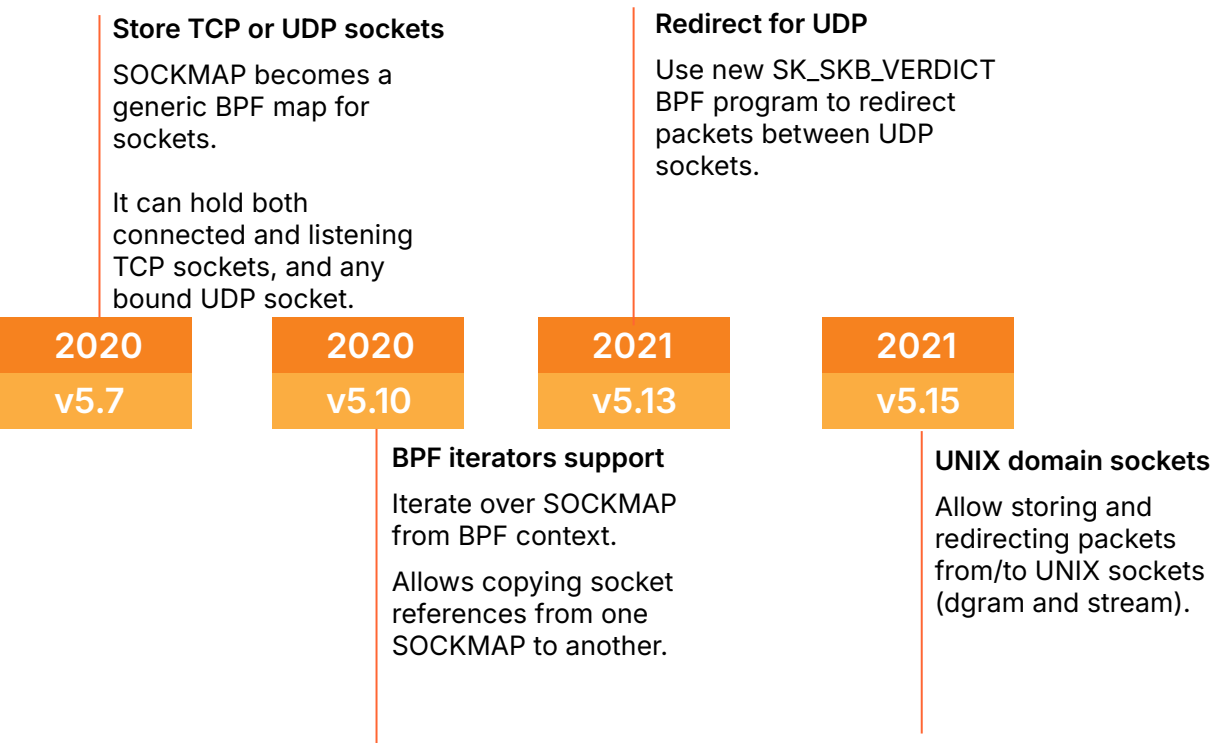
v5.13

BPF iterators support

Iterate over SOCKMAP from BPF context.

Allows copying socket references from one SOCKMAP to another.

How did we get here?



How did we get here?

VSOCK domain sockets

Redirecting from / to
VSOCK sockets –
stream and seqpacket.

2023

v6.4

How did we get here?

VSOCK domain sockets

Redirecting from / to
VSOCK sockets –
stream and seqpacket.

2023

v6.4

202x

v6.x

Your contribution here

The code continues to evolve...

How did we get here?

A third of a century. And it **still** isn't ready. I really need to get my sh*t together..

Linus

How to set up SOCKMAP?




① Open a connected (established) socket

SOCKET



SOCKMAP



weak ref

attached



BPF PROG


active open

```
socket(AF_INET, SOCK_STREAM, IPPROTO_IP) = 8
connect(8, {sa_family=AF_INET,
          sin_port=htons(41895),
          sin_addr=inet_addr("127.0.0.1")}, 16) = 0
```

passive open

```
socket(AF_INET, SOCK_STREAM, IPPROTO_IP) = 7
bind(7, {sa_family=AF_INET,
        sin_port=htons(41895),
        sin_addr=inet_addr("127.0.0.1")}, 16) = 0
listen(7, 4096) = 0
accept(7, NULL, NULL) = 9
```

What sockets can you use?




connected (established) socket:

- ❑ TCP
- ❑ UDP
- ❑ UNIX (STREAM, DGRAM)
- ❑ VSOCK (STREAM, SEQPACKET)

`man 7 {tcp,udp,unix,vsock}`


② Create a BPF map - SOCKMAP or SOCKHASH



use the `bpf()` syscall
or a library wrapper (`ebpf-go`, `libbpf`)

```
bpf(BPF_MAP_CREATE, {map_type=BPF_MAP_TYPE_SOCKMAP,  
                      key_size=4,  
                      value_size=8,  
                      max_entries=1,  
                      map_flags=0,  
                      ...}, 72) = 5
```

② Create a BPF map




or use **bpftool map create** command

```
bpftool map create \
  /sys/fs/bpf/sockmap `# path on bpffs` \
  type sockmap `# sockmap or sockhash` \
  key 4 `# always 4 bytes for sockmap` \
  value 8 `# use 8 bytes for dump to work` \
  entries 1 \
  name sockmap
```

```
bpftool map show pinned /sys/fs/bpf/sockmap
3: sockmap name sockmap flags 0x0
    key 4B value 8B max_entries 1 memlock 328B
```

What BPF maps can you use?




Map types:

- ❑ `BPF_MAP_TYPE_SOCKMAP`
 - ❑ **32-bit integer key**
- ❑ `BPF_MAP_TYPE_SOCKHASH`
 - ❑ **binary blob key**

Not to be confused with
`BPF_MAP_TYPE_REUSEPORT_SOCKARRAY`

③ Load a BPF program




```
bpf(BPF_PROG_LOAD, {prog_type=BPF_PROG_TYPE_SK_MSG,  
                    insn_cnt=6,  
                    insns=0xcf2a70,  
                    expected_attach_type=BPF_SK_MSG_VERDICT,  
                    ...}, 128) = 6
```

Program types:

- ❑ BPF_PROG_TYPE_SK_MSG
- ❑ BPF_PROG_TYPE_SK_SKB


③ Load a BPF program - it uses SOCKMAP



```
# bpftool prog dump xlated id 42
int prog_msg_redir_ingress(struct sk_msg_md * msg):
    0: (18) r2 = map[id:17]
    ...
    5: (95) exit
# bpftool map show id 17
17: sockmap name output flags 0x0
    key 4B value 8B max_entries 1 memlock 328B
    pids sockmap-redir-m(331)
```

④ Attach BPF program to SOCKMAP

SOCKET




SOCKMAP



weak ref

attached




BPF PROG

uses

```
bpf(BPF_PROG_ATTACH, {target_fd=5,  
    attach_bpf_fd=6,  
    attach_type=BPF_SK_MSG_VERDICT,  
    attach_flags=0,  
    replace_bpf_fd=0}, 20) = 0
```

⑤ Insert socket into SOCKMAP



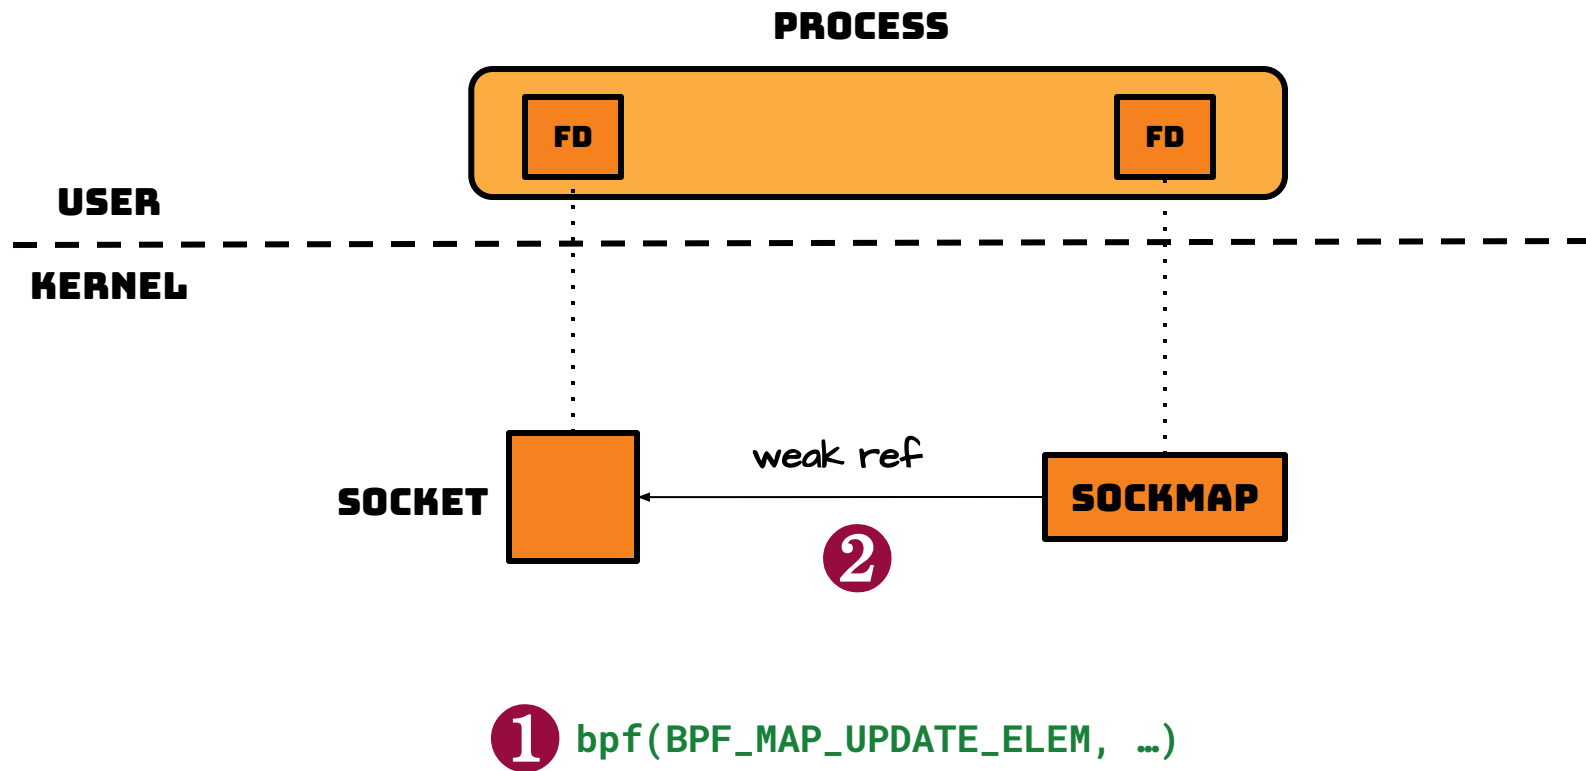
```
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=5,  
    key=0x7ffeb3803870,  
    value=0x7ffeb3803868,  
    flags=BPF_NOEXIST}, 32) = 0
```

⚠ must be done after attaching the program

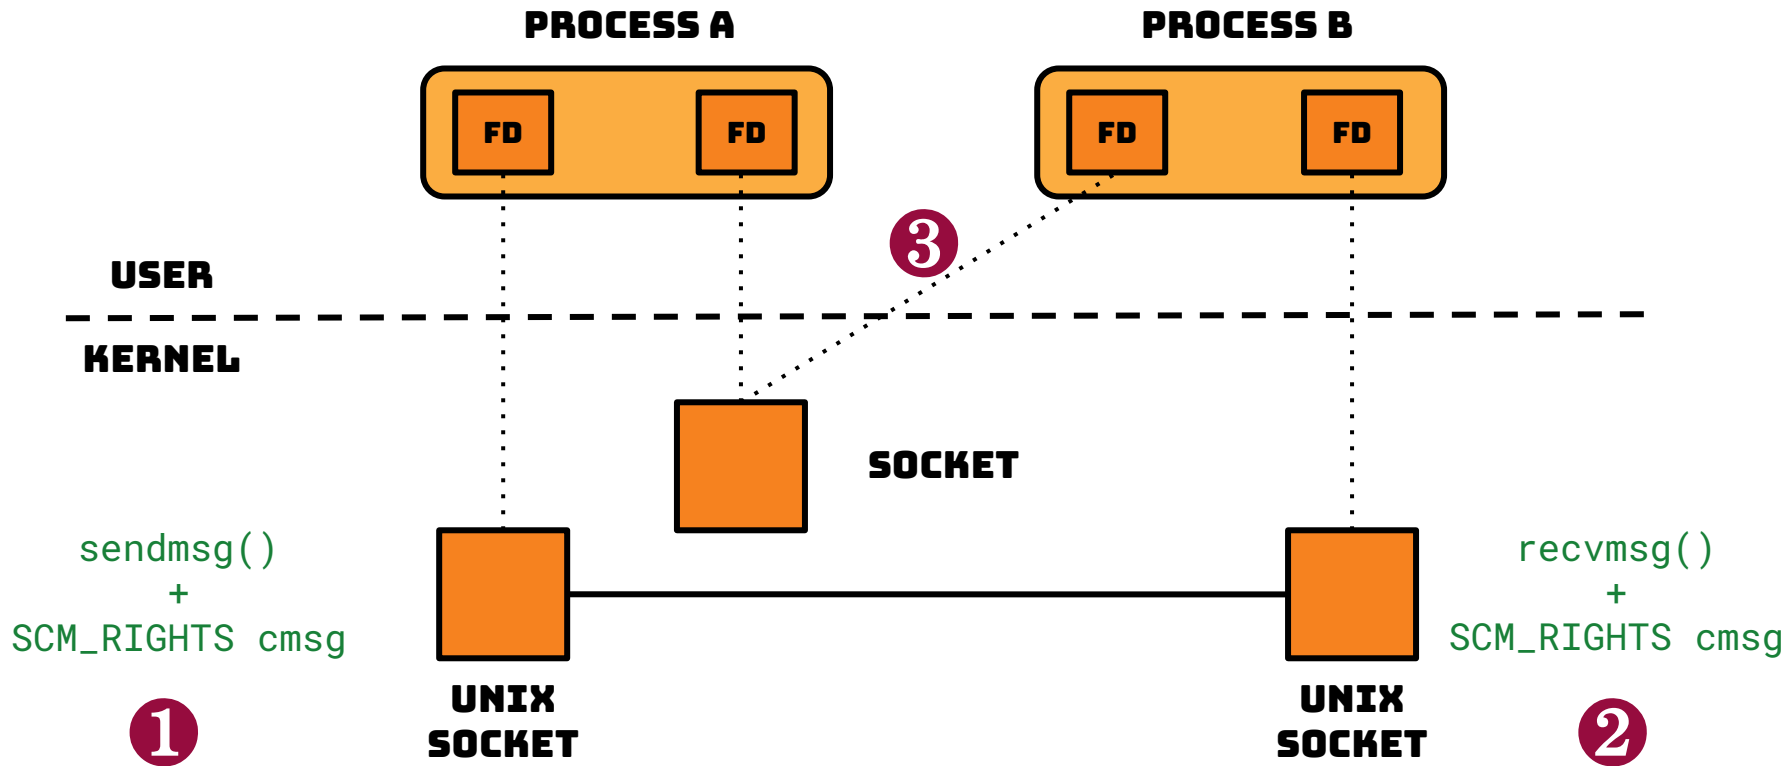
How to get sockets into a SOCKMAP?



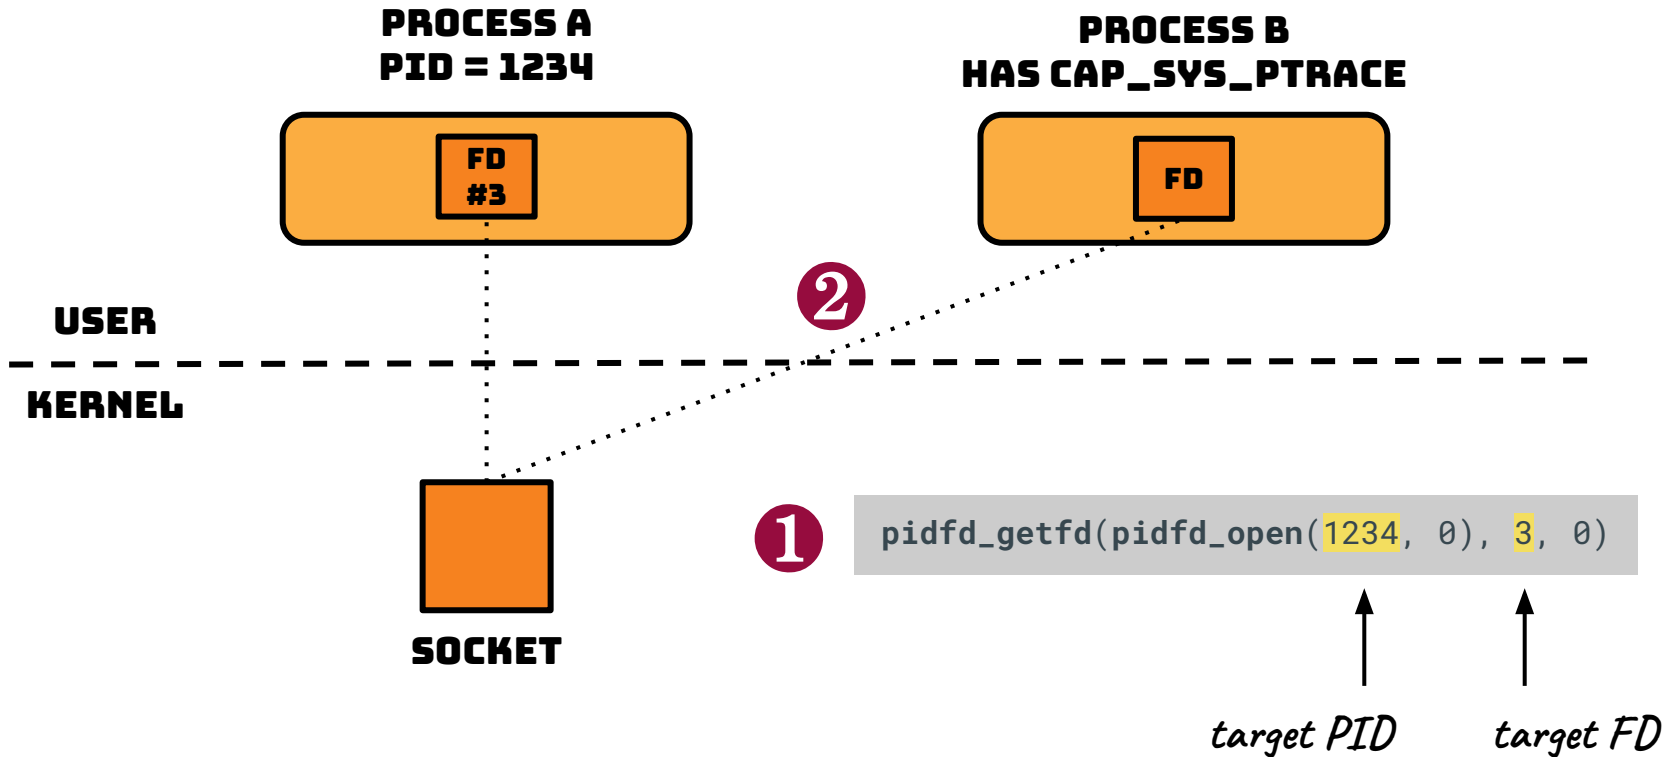
Easy case - Single process



Socket FD handover with SCM_RIGHTS

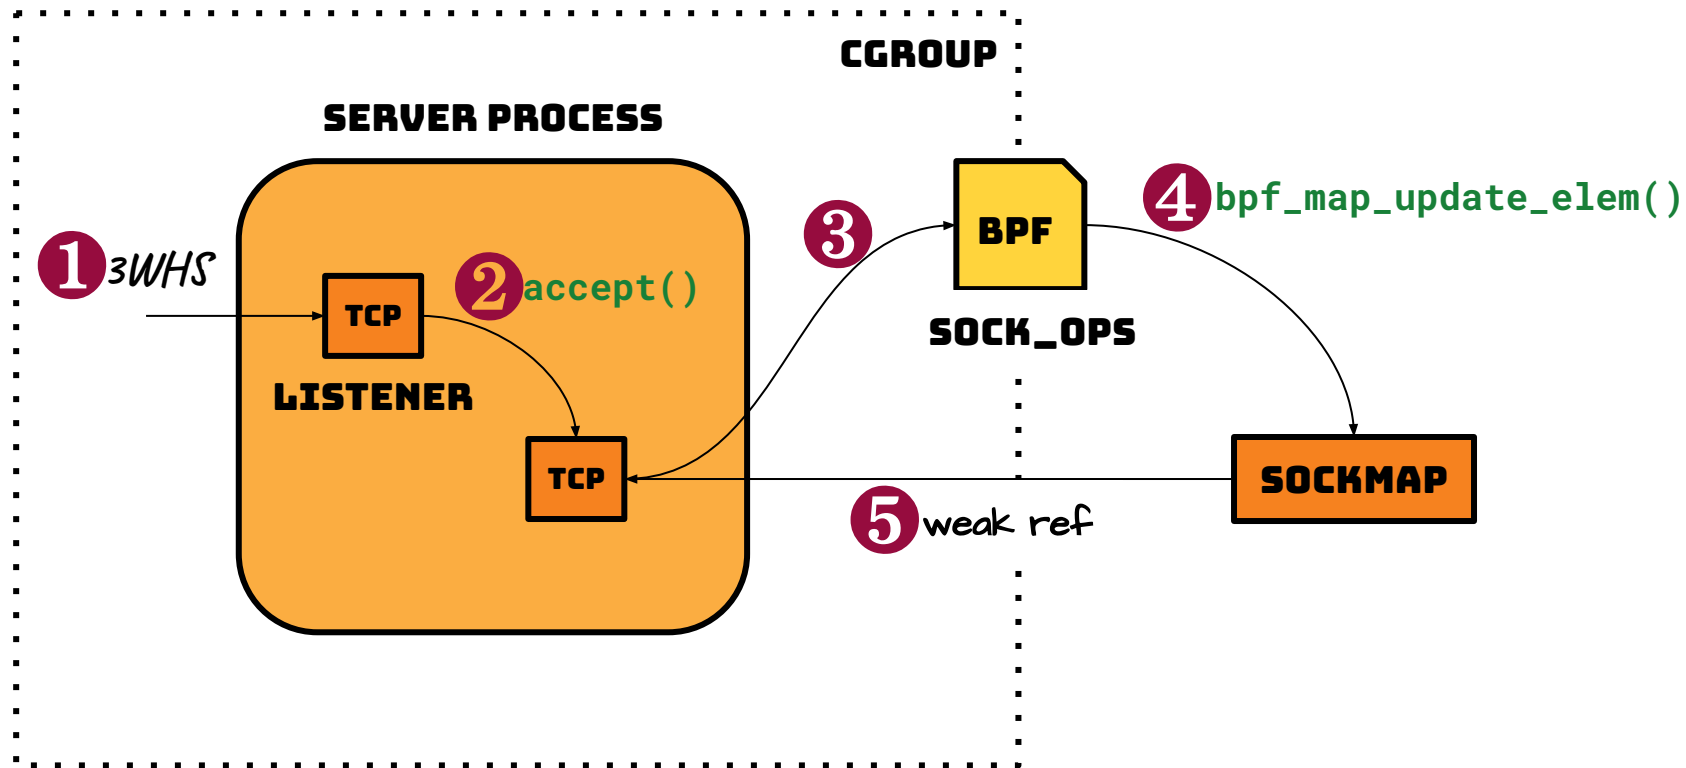


“Steal” a socket FD



Linux 5.6+, requires `CAP_SYS_PTRACE`

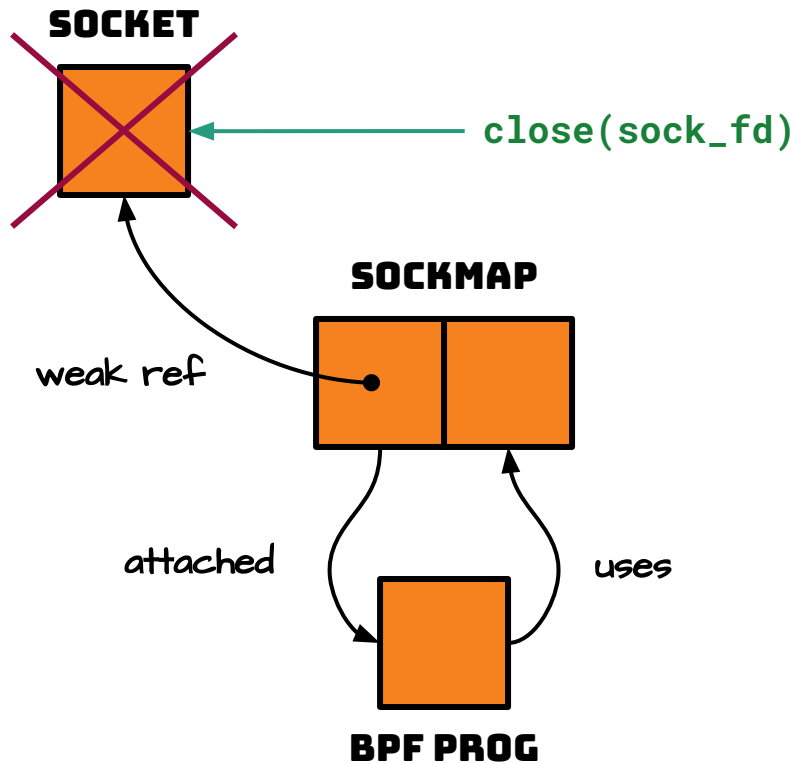
BPF `sock_ops` program attached to cgroup (TCP only)



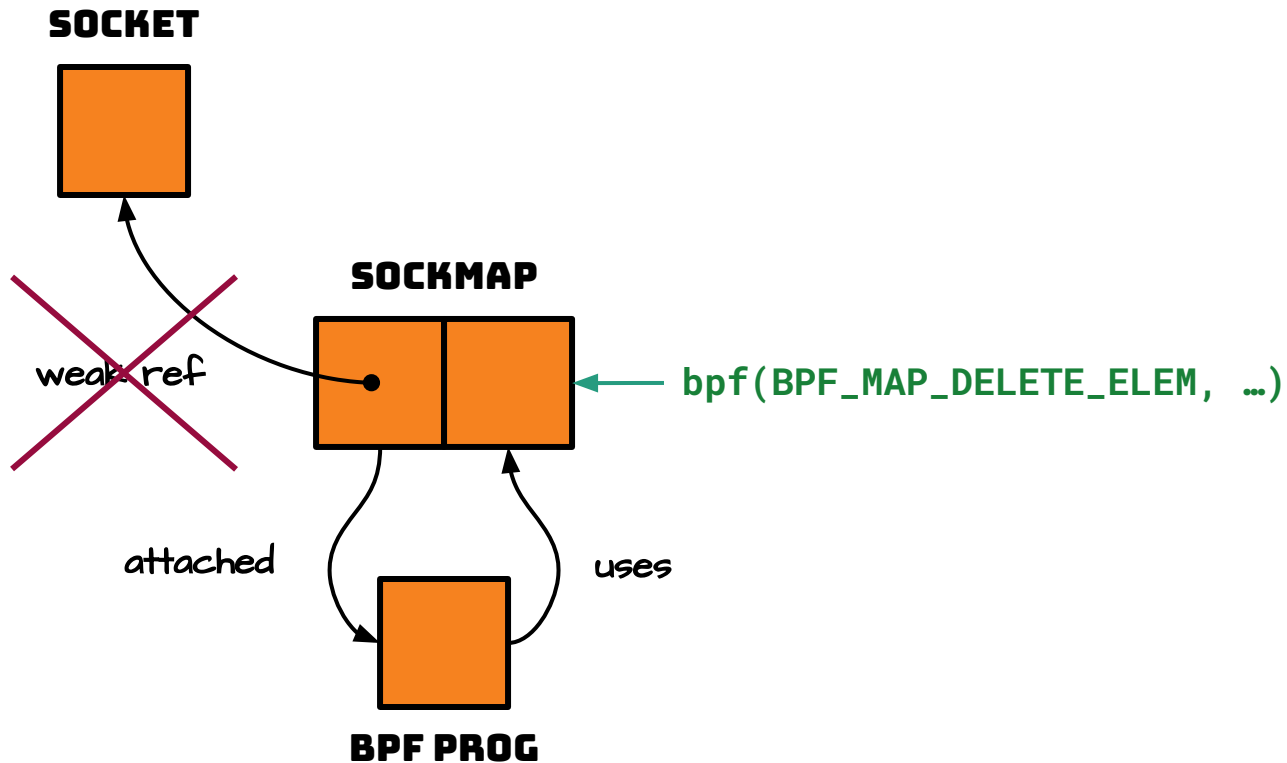
How to tear it down?




(A) destroy the socket



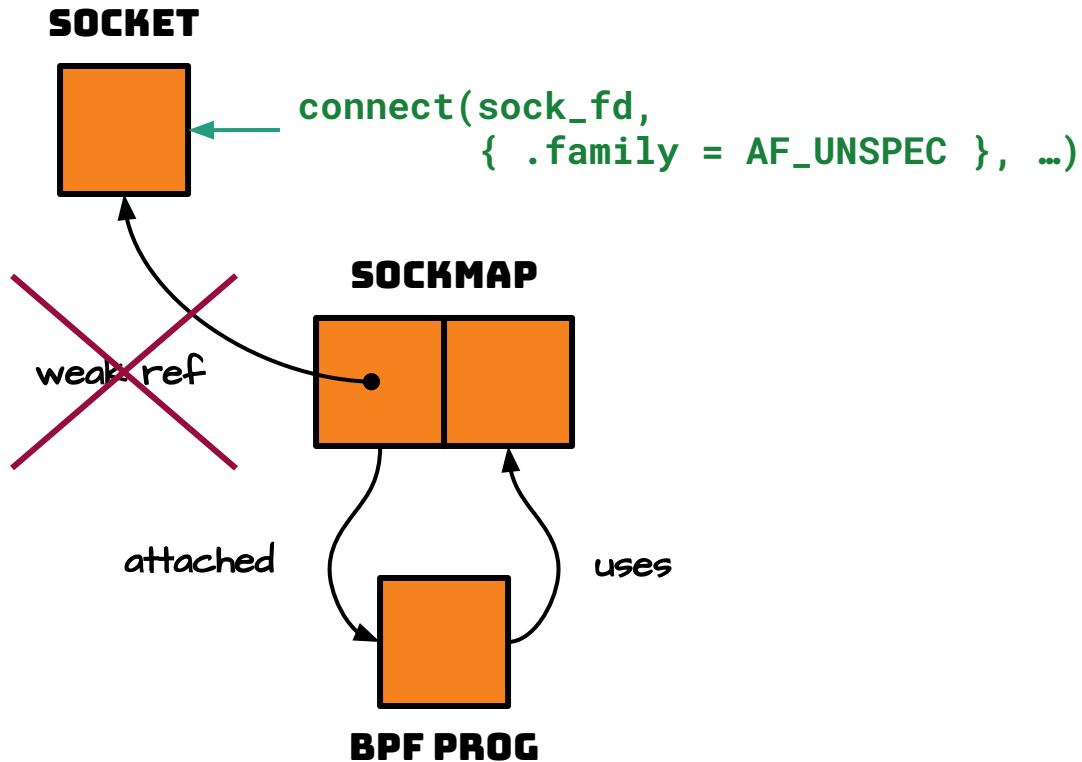
(B) remove socket from sockmap



(C) destroy the sockmap




(D) disconnect the socket (rare)




💡 See `connect(2)` [man page](#)

Supported Socket Splicing Setups





Redirect



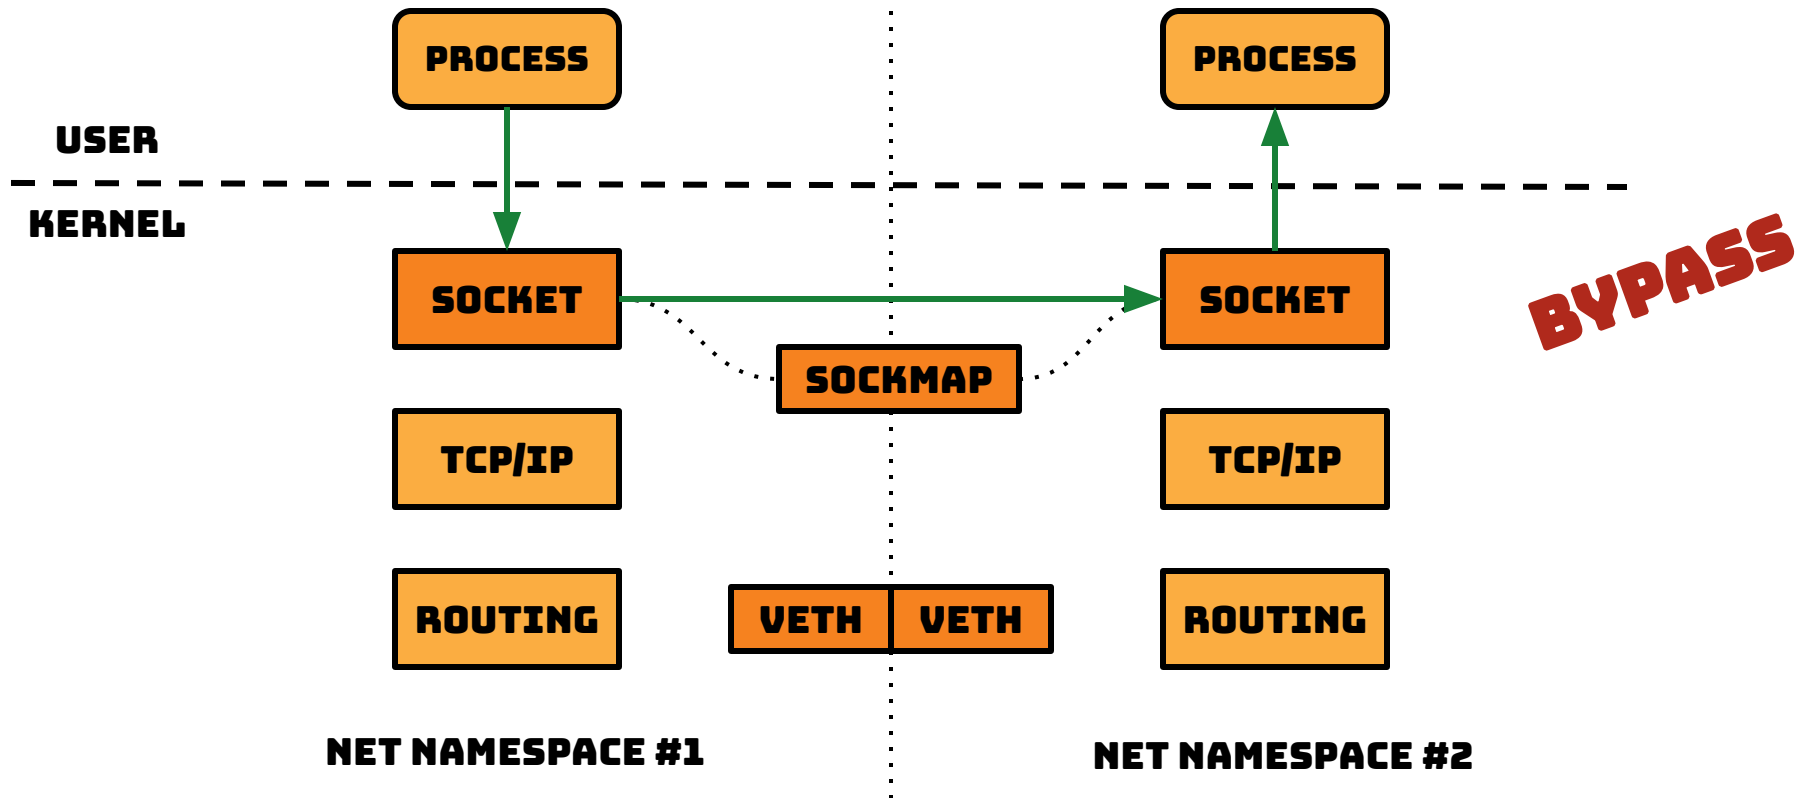
Redirect



send to local




Redirect use case → Bypass for containers





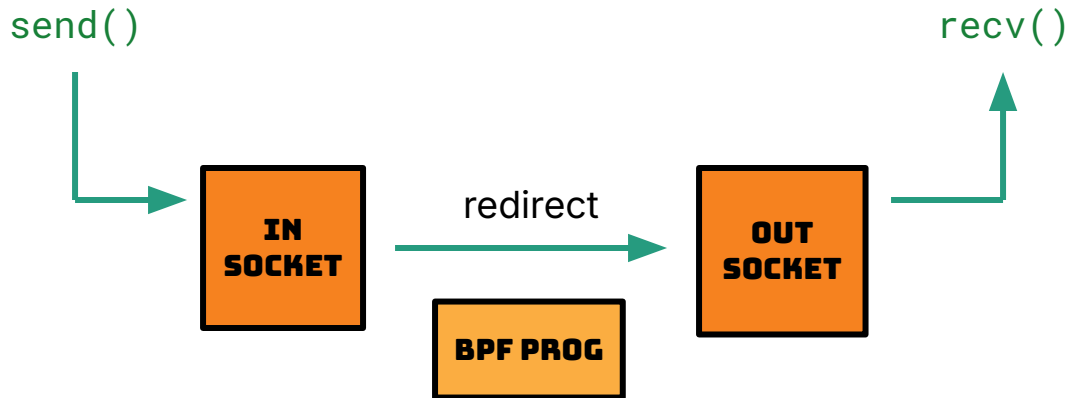
Redirect → send to local



Like `socketpair()` or `pipe()`



Redirect → send to local → How?



BPF_PROG_TYPE_SK_MSG program

→ attached to **BPF_SK_MSG_VERDICT** hook

→ calls **bpf_msg_redirect_hash/map()** with **BPF_F_INGRESS** flag

→ returns **SK_PASS**

selects target socket



Redirect → send to local → Example

```
SEC("sk_msg")
int sk_msg_redir_ingress(struct sk_msg_md *msg)
{
    __u32 key = 0;

    if (msg->remote_port == bpf_htonl(53))
        key = 1;

    return bpf_msg_redirect_map(msg, &sockmap, key, BPF_F_INGRESS);
}
```



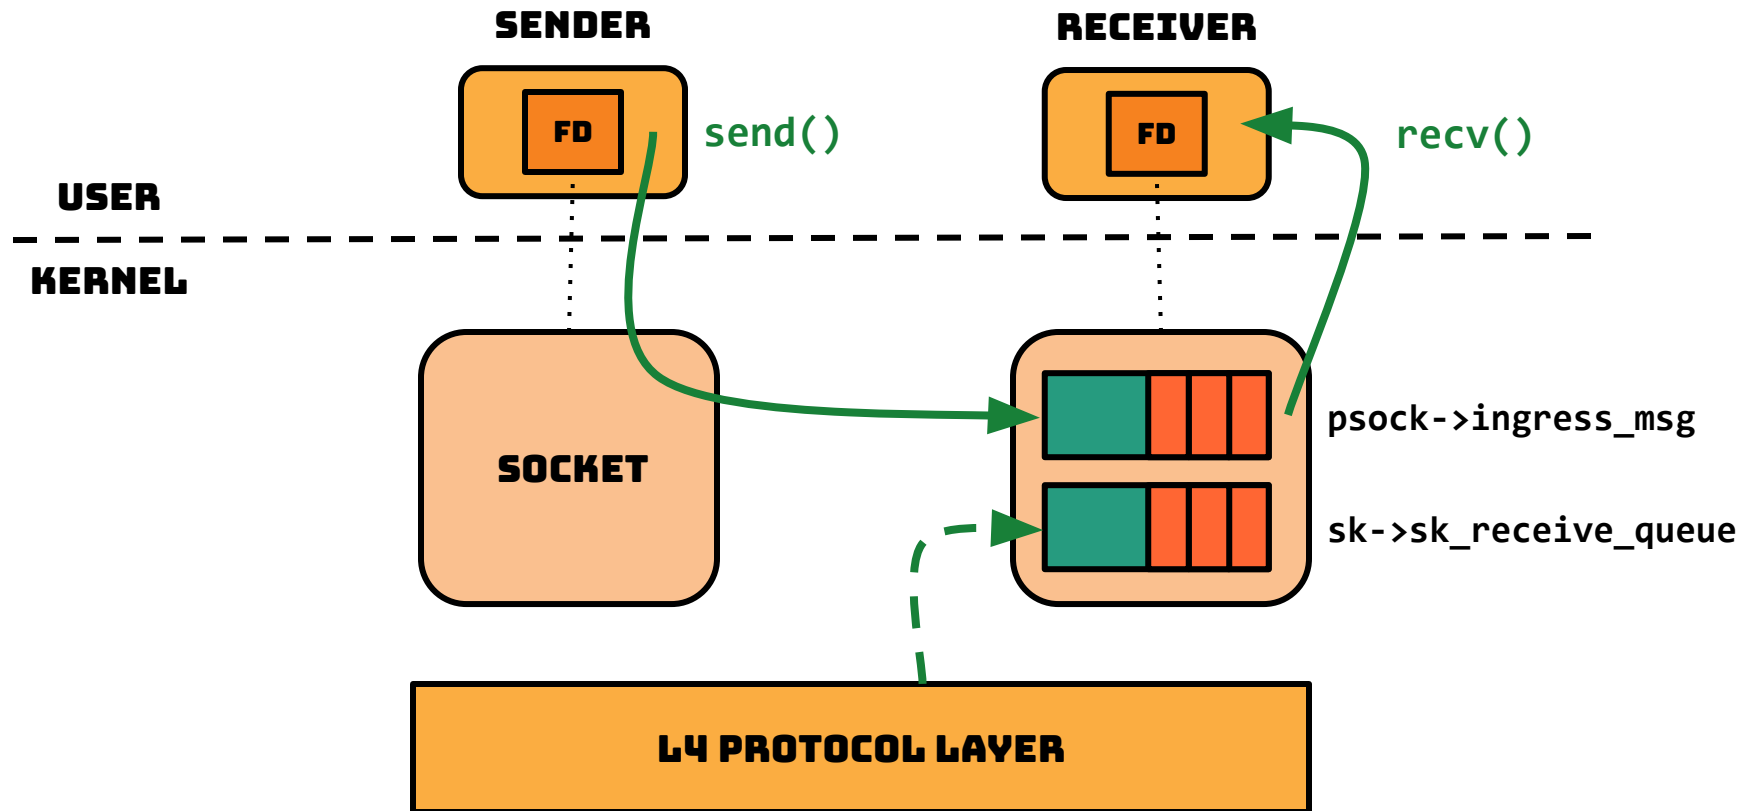

Redirect → send to local → What?

IN → OUT	TCP	UDP	UNIX STR	UNIX DGR	VSOCK STR	VSOCK SEQ
TCP	●	●	●	●	●	●
UDP	●	●	●	●	●	●
UNIX STR	●	●	●	●	●	●
UNIX DGR	●	●	●	●	●	●
VSOCK STR	●	●	●	●	●	●
VSOCK SEQ	●	●	●	●	●	●

TCP to any but VSOCK



Redirect → send to local → Internals



Let's trace what happens in
the sender process during
send-to-local redirect



```
~ # echo 1 > /sys/kernel/tracing/options/funcgraph-retval  
~ # perf ftrace -C 0 -G __sys_sendto --graph-opts noirqs
```

```
sh # perf ftrace -C 0 -G __sys_sendto --graph-opts noirqs
```

```
# tracer: function_graph
```

```
#
```

#	CPU	DURATION	FUNCTION CALLS
#			
0)			__sys_sendto() {
0)			inet_sendmsg() {
0)			tcp_bpf_sendmsg() {
0)	+	23.782 us	sk_msg_alloc();
0)	+	29.818 us	sk_msg_memcpy_from_iter();
0)			sk_psock_msg_verdict() {
0)			bpf_msg_redirect_hash() {
0)		1.764 us	__sock_hash_lookup_elem();
0)		2.599 us	} /* bpf_msg_redirect_hash = 0x1 */
0)		6.048 us	} /* sk_psock_msg_verdict = 0x2 */
0)		1.204 us	sk_msg_return(); /* = 0x1 */
0)			tcp_bpf_sendmsg_redir() {
0)		1.507 us	__sk_mem_schedule(); /* = 0x1 */
0)			sock_def_readable() {
0)			__wake_up_sync_key() {
0)			__wake_up_common() {
0)			pollwake() {
0)			default_wake_function() {
0)	+	12.872 us	try_to_wake_up(); /* = 0x1 */
0)	+	13.207 us	} /* default_wake_function = 0x1 */
0)	+	13.650 us	} /* pollwake = 0x1 */
0)	+	14.583 us	} /* __wake_up_common = 0x1 */
0)	+	16.905 us	} /* __wake_up_sync_key = 0x4b18 */
0)	+	19.696 us	} /* sock_def_readable = 0x0 */
0)	+	63.929 us	} /* tcp_bpf_sendmsg_redir = 0x0 */
0)	!	168.534 us	} /* tcp_bpf_sendmsg = 0x7 */
0)	!	172.112 us	} /* inet_sendmsg = 0x7 */
0)	!	194.533 us	} /* __sys_sendto = 0x7 */

*irrelevant
bits
omitted*

```
sh # perf ftrace -C 0 -G __sys_sendto --graph-opts noirqs
```

```
# tracer: function_graph
```

```
#
```

```
# CPU    DURATION
```

```
FUNCTION CALLS
```

```
# |      |      |      |      |      |      |
```

```
0)
```

```
__sys_sendto() {
```

```
0)
```

```
inet_sendmsg() {
```

```
0)
```

```
tcp_bpf_sendmsg() {
```

```
0) + 23.782 us
```

```
sk_msg_alloc();
```

```
0) + 29.818 us
```

```
sk_msg_memcpy_from_iter();
```

```
0)
```

```
sk_psock_msg_verdict() {
```

```
0)
```

```
bpf_msg_redirect_hash() {
```

```
0) 1.764 us
```

```
__sock_hash_lookup_elem();
```

```
0) 2.599 us
```

```
} /* bpf_msg_redirect_hash = 0x1 */
```

```
0) 6.048 us
```

```
} /* sk_psock_msg_verdict = 0x2 */
```

```
0) 1.204 us
```

```
sk_msg_return(); /* = 0x1 */
```

```
0)
```

```
tcp_bpf_sendmsg_redir() {
```

```
0) 1.507 us
```

```
__sk_mem_schedule(); /* = 0x1 */
```

```
0)
```

```
sock_def_readable() {
```

```
0)
```

```
__wake_up_sync_key() {
```

```
0)
```

```
__wake_up_common() {
```

```
0)
```

```
pollwake() {
```

```
0)
```

```
default_wake_function() {
```

```
0) + 12.872 us
```

```
try_to_wake_up(); /* = 0x1 */
```

```
0) + 13.207 us
```

```
} /* default_wake_function = 0x1 */
```

```
0) + 13.650 us
```

```
} /* pollwake = 0x1 */
```

```
0) + 14.583 us
```

```
} /* __wake_up_common = 0x1 */
```

```
0) + 16.905 us
```

```
} /* __wake_up_sync_key = 0x4b18 */
```

```
0) + 19.696 us
```

```
} /* sock_def_readable = 0x0 */
```

```
0) + 63.929 us
```

```
} /* tcp_bpf_sendmsg_redir = 0x0 */
```

```
0) ! 168.534 us
```

```
} /* tcp_bpf_sendmsg = 0x7 */
```

```
0) ! 172.112 us
```

```
} /* inet_sendmsg = 0x7 */
```

```
0) ! 194.533 us
```

```
} /* __sys_sendto = 0x7 */
```

Entry to syscall,
allocate & initialize
the message

```
sh # perf ftrace -C 0 -G __sys_sendto --graph-opts noirqs
```

```
# tracer: function_graph
```

```
#
```

```
# CPU DURATION
```

```
FUNCTION CALLS
```

#							
0)				__sys_sendto()	{		
0)				inet_sendmsg()	{		
0)				tcp_bpf_sendmsg()	{		
0)	+	23.782	us	sk_msg_alloc()			
0)	+	29.818	us	sk_msg_memcpy_from_iter()			
0)				sk_psock_msg_verdict()	{		
0)				bpf_msg_redirect_hash()	{		
0)		1.764	us	__sock_hash_lookup_elem()			
0)		2.599	us	}	/* bpf_msg_redirect_hash = 0x1 */		
0)		6.048	us	}	/* sk_psock_msg_verdict = 0x2 */		
0)		1.204	us	sk_msg_return()	/* = 0x1 */		
0)				tcp_bpf_sendmsg_redir()	{		
0)		1.507	us	__sk_mem_schedule()	/* = 0x1 */		
0)				sock_def_readable()	{		
0)				__wake_up_sync_key()	{		
0)				__wake_up_common()	{		
0)				pollwake()	{		
0)				default_wake_function()	{		
0)	+	12.872	us	try_to_wake_up()	/* = 0x1 */		
0)	+	13.207	us	}	/* default_wake_function = 0x1 */		
0)	+	13.650	us	}	/* pollwake = 0x1 */		
0)	+	14.583	us	}	/* __wake_up_common = 0x1 */		
0)	+	16.905	us	}	/* __wake_up_sync_key = 0x4b18 */		
0)	+	19.696	us	}	/* sock_def_readable = 0x0 */		
0)	+	63.929	us	}	/* tcp_bpf_sendmsg_redir = 0x0 */		
0)	!	168.534	us	}	/* tcp_bpf_sendmsg = 0x7 */		
0)	!	172.112	us	}	/* inet_sendmsg = 0x7 */		
0)	!	194.533	us	}	/* __sys_sendto = 0x7 */		

Run BPF program,
find the destination
socket in the BPF map

```
sh # perf ftrace -C 0 -G __sys_sendto --graph-opts noirqs
```

```
# tracer: function_graph
```

```
#
```

#	CPU	DURATION	FUNCTION CALLS
#			
0)			__sys_sendto() {
0)			inet_sendmsg() {
0)			tcp_bpf_sendmsg() {
0)	+	23.782 us	sk_msg_alloc();
0)	+	29.818 us	sk_msg_memcpy_from_iter();
0)			sk_psock_msg_verdict() {
0)			bpf_msg_redirect_hash() {
0)		1.764 us	__sock_hash_lookup_elem();
0)		2.599 us	} /* bpf_msg_redirect_hash = 0x1 */
0)		6.048 us	} /* sk_psock_msg_verdict = 0x2 */
0)		1.204 us	sk_msg_return(); /* = 0x1 */
0)			tcp_bpf_sendmsg_redir() {
0)		1.507 us	__sk_mem_schedule(); /* = 0x1 */
0)			sock_def_readable() {
0)			__wake_up_sync_key() {
0)			__wake_up_common() {
0)			pollwake() {
0)			default_wake_function() {
0)	+	12.872 us	try_to_wake_up(); /* = 0x1 */
0)	+	13.207 us	} /* default_wake_function = 0x1 */
0)	+	13.650 us	} /* pollwake = 0x1 */
0)	+	14.583 us	} /* __wake_up_common = 0x1 */
0)	+	16.905 us	} /* __wake_up_sync_key = 0x4b18 */
0)	+	19.696 us	} /* sock_def_readable = 0x0 */
0)	+	63.929 us	} /* tcp_bpf_sendmsg_redir = 0x0 */
0)	!	168.534 us	} /* tcp_bpf_sendmsg = 0x7 */
0)	!	172.112 us	} /* inet_sendmsg = 0x7 */
0)	!	194.533 us	} /* __sys_sendto = 0x7 */

Deliver message to receiver's queue, and notify the receiver

```
sh # perf ftrace -C 0 -G __sys_sendto --graph-opts noirqs
```

```
# tracer: function_graph
```

```
#
```

#	CPU	DURATION	FUNCTION CALLS
#			
0)			__sys_sendto() {
0)			inet_sendmsg() {
0)			tcp_bpf_sendmsg() {
0)	+	23.782 us	sk_msg_alloc();
0)	+	29.818 us	sk_msg_memcpy_from_iter();
0)			sk_psock_msg_verdict() {
0)			bpf_msg_redirect_hash() {
0)		1.764 us	__sock_hash_lookup_elem();
0)		2.599 us	} /* bpf_msg_redirect_hash = 0x1 */
0)		6.048 us	} /* sk_psock_msg_verdict = 0x2 */
0)		1.204 us	sk_msg_return(); /* = 0x1 */
0)			tcp_bpf_sendmsg_redir() {
0)		1.507 us	__sk_mem_schedule(); /* = 0x1 */
0)			sock_def_readable() {
0)			__wake_up_sync_key() {
0)			__wake_up_common() {
0)			pollwake() {
0)			default_wake_function() {
0)	+	12.872 us	try_to_wake_up(); /* = 0x1 */
0)	+	13.207 us	} /* default_wake_function = 0x1 */
0)	+	13.650 us	} /* pollwake = 0x1 */
0)	+	14.583 us	} /* __wake_up_common = 0x1 */
0)	+	16.905 us	} /* __wake_up_sync_key = 0x4b18 */
0)	+	19.696 us	} /* sock_def_readable = 0x0 */
0)	+	63.929 us	} /* tcp_bpf_sendmsg_redir = 0x0 */
0)	!	168.534 us	} /* tcp_bpf_sendmsg = 0x7 */
0)	!	172.112 us	} /* inet_sendmsg = 0x7 */
0)	!	194.533 us	} /* __sys_sendto = 0x7 */

Wake up the receiver


Redirect



send to egress



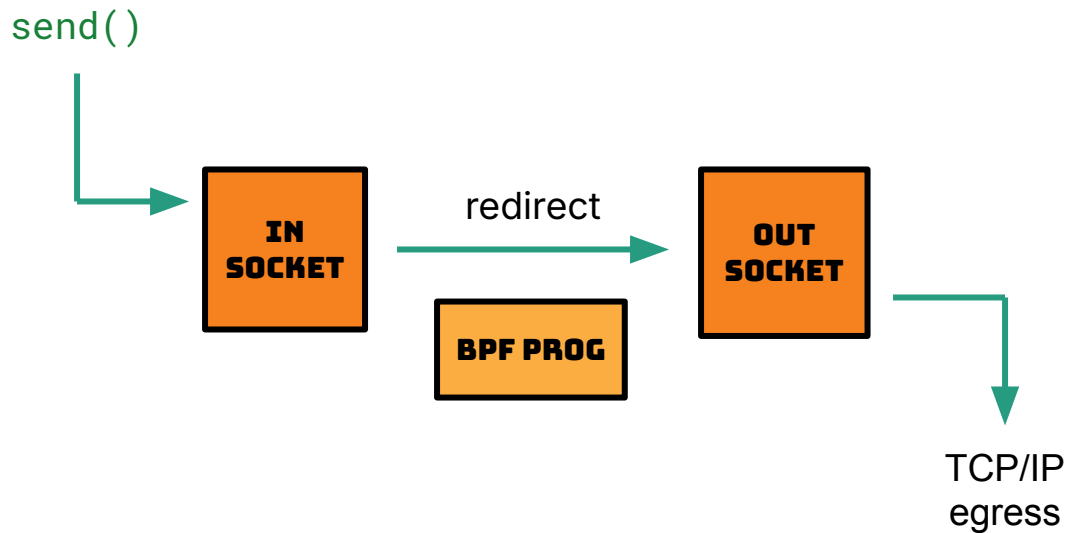
Redirect → send to egress



Sort of like `splice()`
(pipe → socket)



Redirect → send to egress → How?



`BPF_PROG_TYPE_SK_MSG` prog

→ attached to `BPF_SK_MSG_VERDICT` hook

→ calls `bpf_msg_redirect_hash/map()` without any flags

→ returns `SK_PASS`




Redirect → send to egress → What?

IN → OUT	TCP	UDP	UNIX STR	UNIX DGR	VSOCK STR	VSOCK SEQ
TCP	●	●	●	●	●	●
UDP	●	●	●	●	●	●
UNIX STR	●	●	●	●	●	●
UNIX DGR	●	●	●	●	●	●
VSOCK STR	●	●	●	●	●	●
VSOCK SEQ	●	●	●	●	●	●

TCP to TCP only

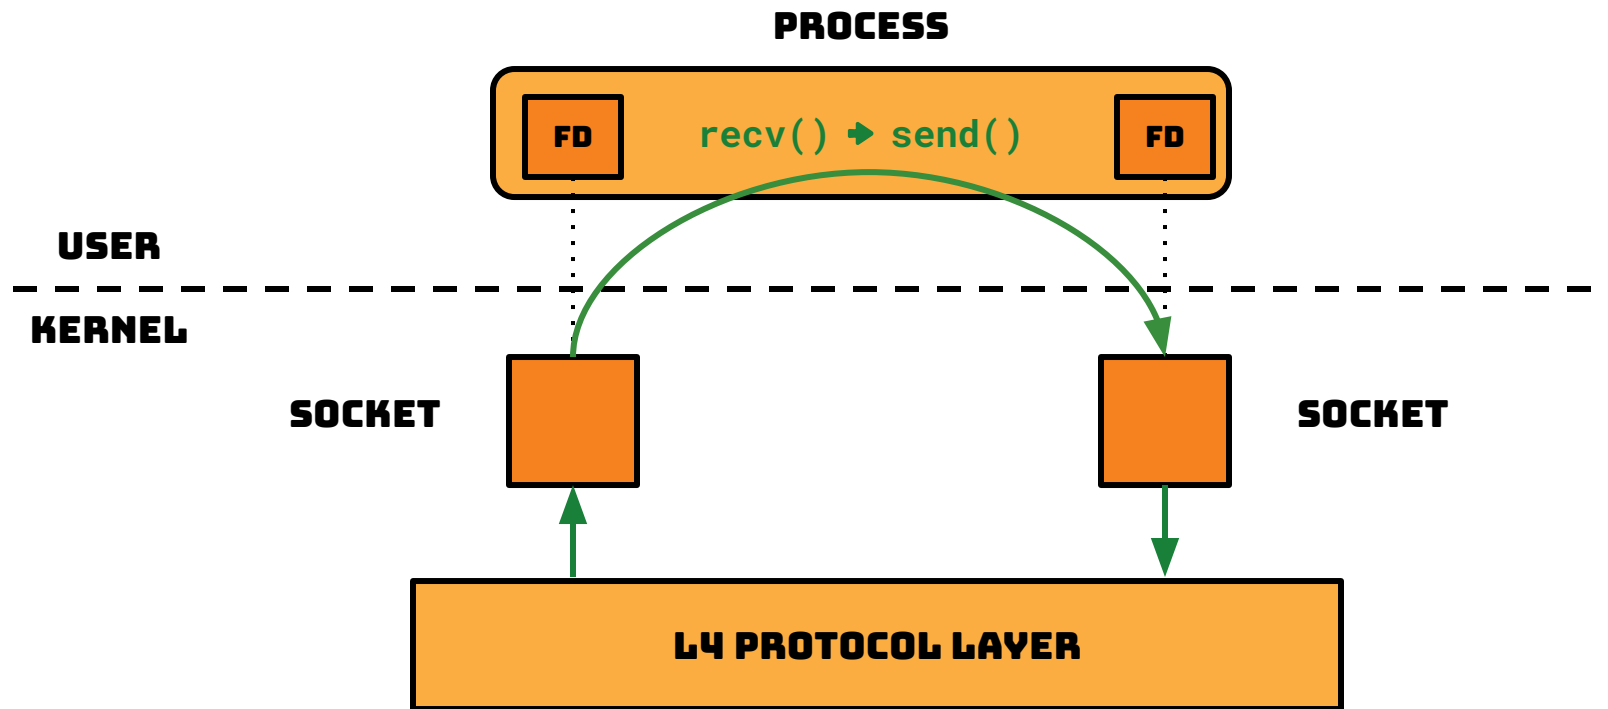
Redirect



ingress to egress



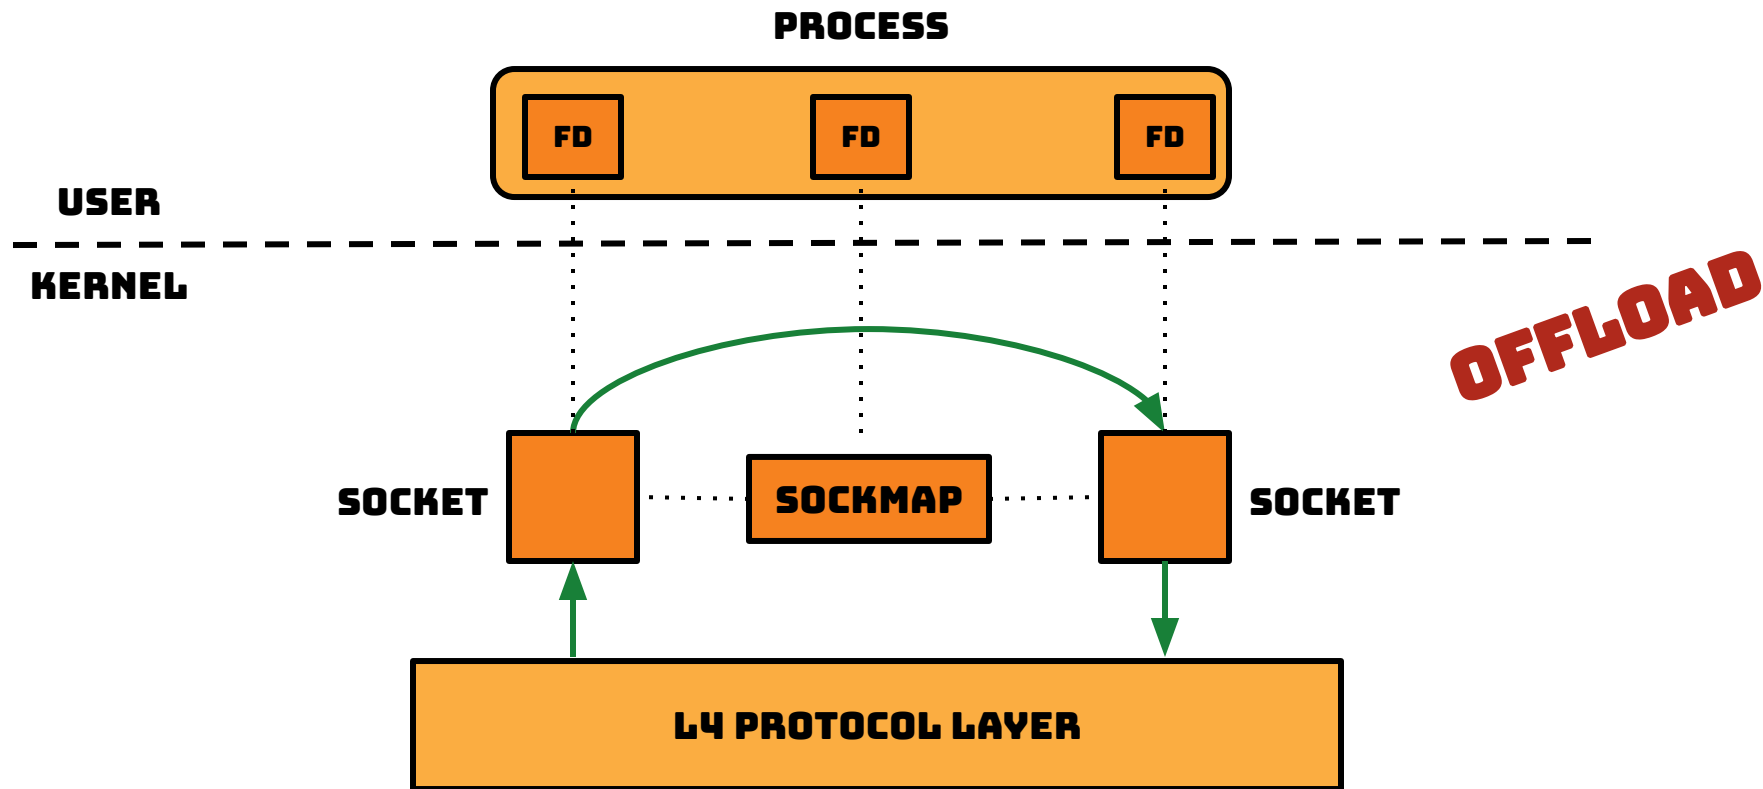
Redirect use case → L7 network proxy



Examples: `socat`, `systemd-socket-proxyd`




Redirect use case → L7 network proxy






Redirect → ingress to egress



Like double `splice()`
(socket → pipe → socket)



Redirect → ingress to egress → How?



`BPF_PROG_TYPE_SK_SKB` prog

→ attached to `BPF_SK_SKB_VERDICT` hook

→ calls `bpf_sk_redirect_hash/map()` without any flags

→ returns `SK_PASS`



Redirect → ingress to egress → What?

IN → OUT	TCP	UDP	UNIX STR	UNIX DGR	VSOCK STR	VSOCK SEQ
TCP	●	●	●	●	●	●
UDP	●	●	●	●	●	●
UNIX STR	●	●	●	●	●	●
UNIX DGR	●	●	●	●	●	●
VSOCK STR	●	●	●	●	●	●
VSOCK SEQ	●	●	●	●	●	●

any to any



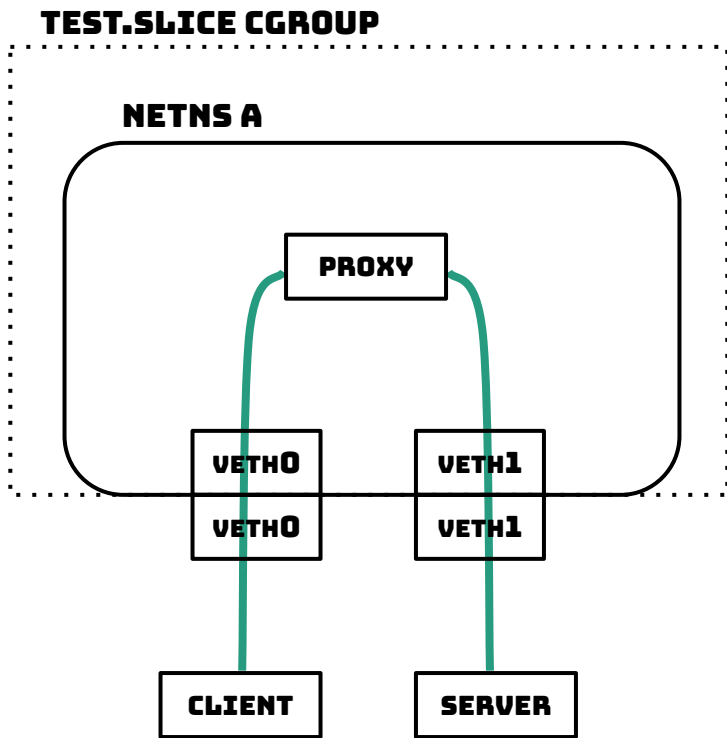
Another benchmark time!



Image credit: [Stopwatch Timer Chronometer](#) - Free photo on Pixabay



Ping-pong test through a TCP proxy





Ping-pong test through a TCP proxy

Spawn proxy inside the test.slice cgroup and in netns A

```
# (  
    echo $BASHPID > /sys/fs/cgroup/test.slice/cgroup.procs  
    taskset -c 0 ip netns exec A \  
        ./tcp_proxy -proxy="10.100.0.10:1111" -target="10.200.0.1:2222" &  
)
```

Start TCP server in main netns

```
# taskset -c 2 sockperf server -i 10.200.0.1 -p 2222 --tcp &
```

Run TCP client in main netns

```
# taskset -c 4 sockperf ping-pong -i 10.100.0.10 -p 1111 --tcp --time 30
```



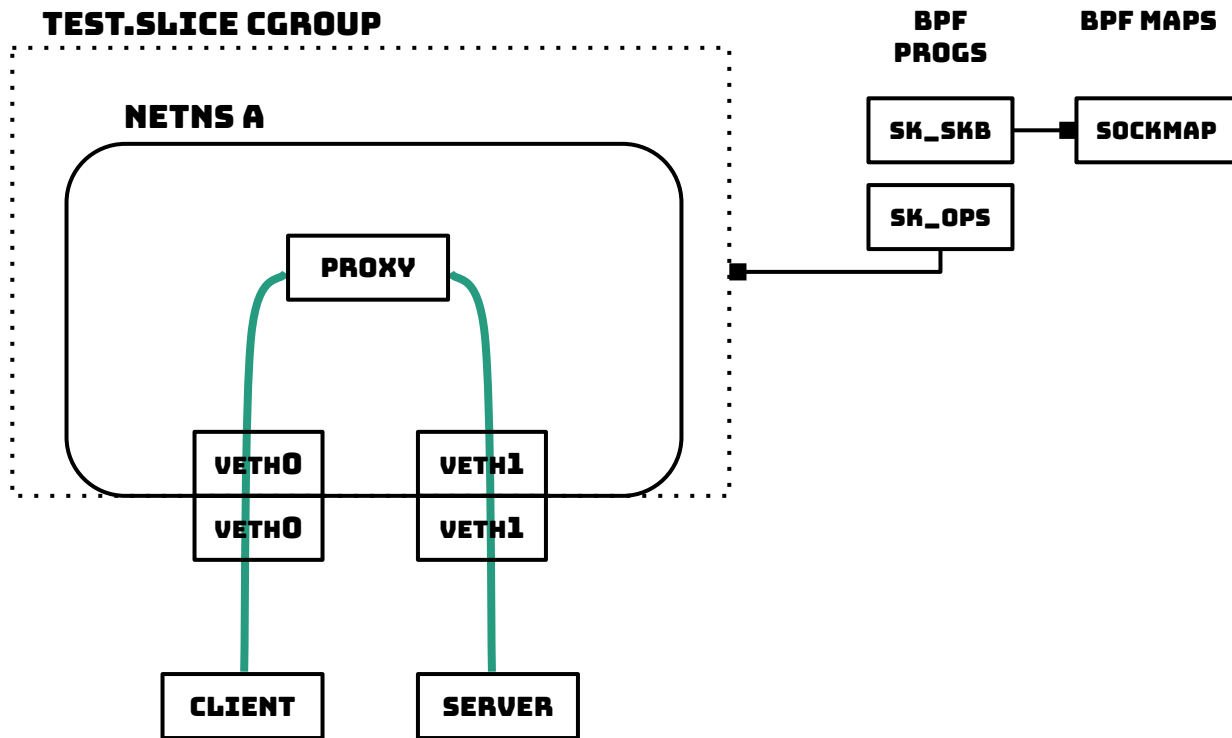
Ping-pong test through a TCP proxy

```
# taskset -c 4 sockperf ping-pong -i 10.100.0.10 -p 1111 --tcp --time 30
sockperf: Starting test...
sockperf: Test end (interrupted by timer)
sockperf: Test ended
sockperf: [Total Run] RunTime=30.001 sec; Warm up time=400 msec; ...
sockperf: ===== Printing statistics for Server No: 0
sockperf: [Valid Duration] RunTime=29.550 sec; SentMessages=619300; ReceivedMessages=619300
sockperf: =====> avg-latency=23.827 (std-dev=16.129)
sockperf: # dropped messages = 0; # duplicated messages = 0; # out-of-order messages = 0
sockperf: Summary: Latency is 23.827 usec
sockperf: Total 619300 observations; each percentile contains 6193.00 observations
sockperf: ---> <MAX> observation = 1449.742
sockperf: ---> percentile 99.999 = 1382.814
sockperf: ---> percentile 99.990 = 1034.076
sockperf: ---> percentile 99.900 = 82.949
sockperf: ---> percentile 99.000 = 33.360
sockperf: ---> percentile 90.000 = 25.817
sockperf: ---> percentile 75.000 = 24.732
sockperf: ---> percentile 50.000 = 23.810
sockperf: ---> percentile 25.000 = 22.749
sockperf: ---> <MIN> observation = 15.487
```

24 ± 16 μsec



Same with SOCKMAP bypass




```
SEC("sockops")
int sockops_prog(struct bpf_sock_ops *ctx)
{
    enum conn_dir dir;
    __u64 cookie;

    switch (ctx->op) {
    case BPF SOCK_OPS_ACTIVE_ESTABLISHED_CB:
        dir = OUTGOING;
        break;
    case BPF SOCK_OPS_PASSIVE_ESTABLISHED_CB:
        dir = INCOMING;
        break;
    default:
        goto out;
    }

    cookie = bpf_get_socket_cookie(ctx);
    bpf_sock_map_update(ctx, &sock_map, &dir, /* flags= */ 0);
    bpf_map_update_elem(&conn_map, &cookie, &dir, BPF_ANY);
out:
    return SK_PASS;
}
```


```
SEC("sk_skb")
int redir_skb_prog(struct __sk_buff *skb)
{
    __u64 cookie = bpf_get_socket_cookie(skb);
    enum conn_dir *v, target;

    v = bpf_map_lookup_elem(&conn_map, &cookie);
    if (!v)
        goto err;

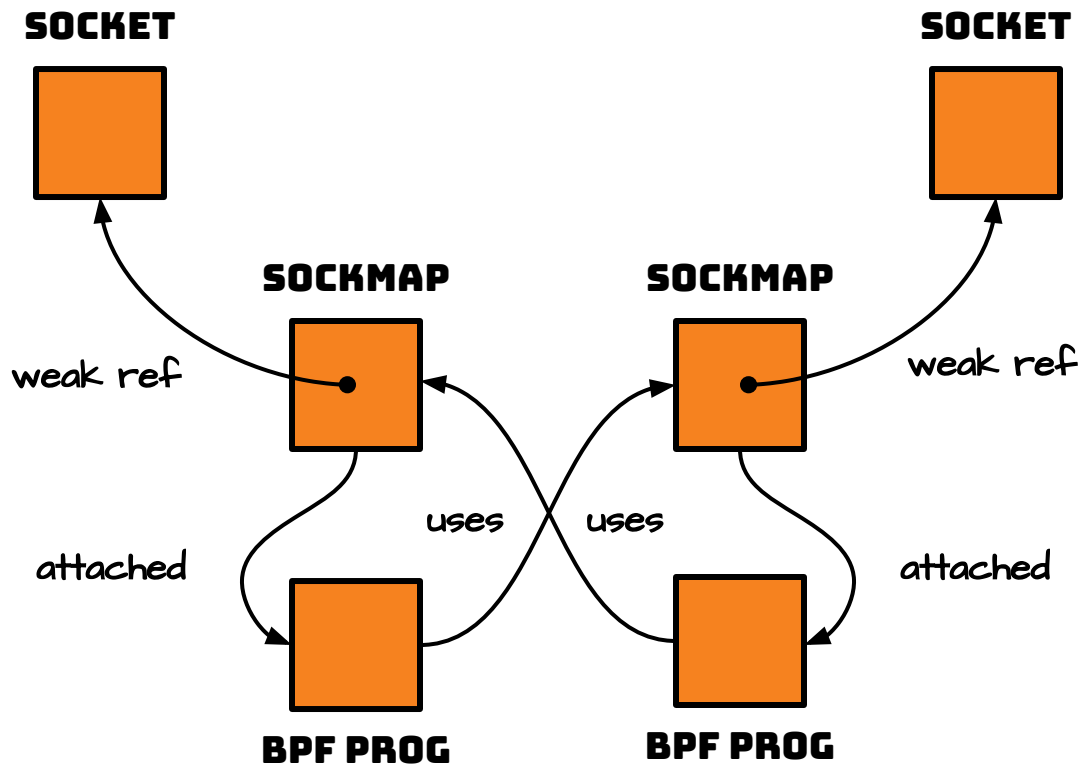
    switch (*v) {
    case INCOMING:
        target = OUTGOING;
        break;
    case OUTGOING:
        target = INCOMING;
        break;
    default:
        goto err;
    }

    return bpf_sk_redirect_map(skb, &sock_map, target, /* flags= */ 0);
err:
    __sync_fetch_and_add(&redir_errors, 1);
    return SK_DROP;
}
```

Two sockets in one sockmap



Two sockets in two sockmaps





Same with SOCKMAP bypass

```
# taskset -c 4 sockperf ping-pong -i 10.100.0.10 -p 1111 --tcp --time 30
sockperf: Starting test...
sockperf: Test end (interrupted by timer)
sockperf: Test ended
sockperf: [Total Run] RunTime=30.000 sec; Warm up time=400 msec; ...
sockperf: ===== Printing statistics for Server No: 0
sockperf: [Valid Duration] RunTime=29.552 sec; SentMessages=738402; ReceivedMessages=738402
sockperf: =====> avg-latency=19.973 (std-dev=5.672)
sockperf: # dropped messages = 0; # duplicated messages = 0; # out-of-order messages = 0
sockperf: Summary: Latency is 19.973 usec
sockperf: Total 738402 observations; each percentile contains 7384.02 observations
sockperf: ---> <MAX> observation = 1238.486
sockperf: ---> percentile 99.999 = 696.232
sockperf: ---> percentile 99.990 = 147.685
sockperf: ---> percentile 99.900 = 64.345
sockperf: ---> percentile 99.000 = 38.314
sockperf: ---> percentile 90.000 = 22.524
sockperf: ---> percentile 75.000 = 21.908
sockperf: ---> percentile 50.000 = 18.355
sockperf: ---> percentile 25.000 = 17.427
sockperf: ---> <MIN> observation = 13.225
```

20 ± 6.0 µsec



Ping-pong test through a TCP proxy with and without SOCKMAP bypass

before: $24 \pm 16 \mu\text{sec}$

↓ - 17%

after: $20 \pm 6.0 \mu\text{sec}$

Run the benchmark yourself:

<https://github.com/jsitnicki/sockmap-project/tree/main/examples/ingress-to-egress>



But...

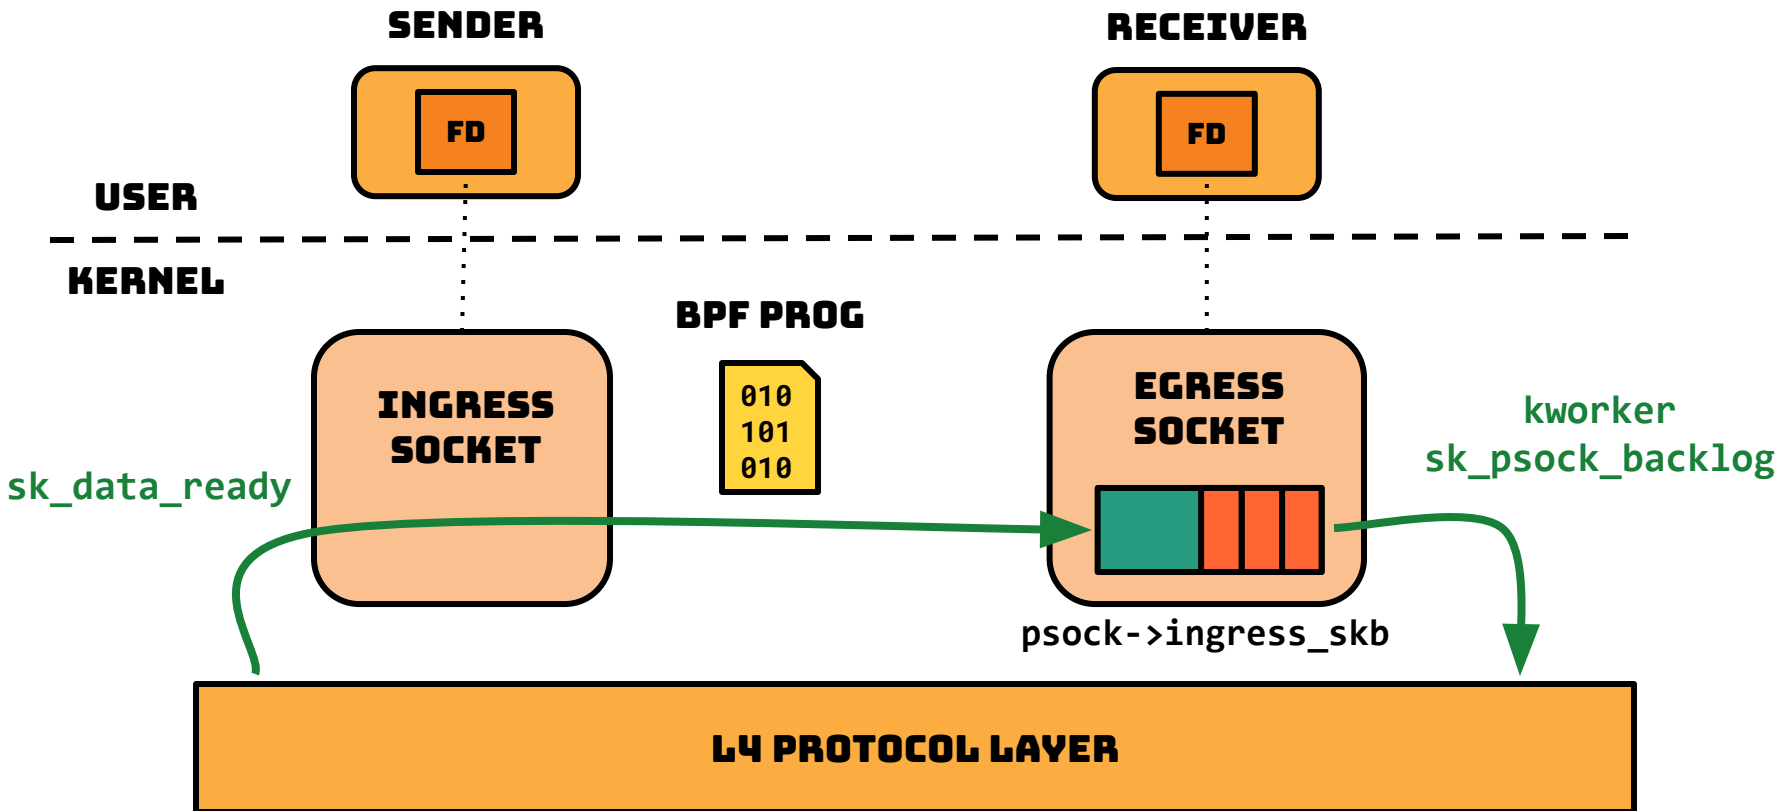


“there is always a but in this imperfect world!”

— Anne Brontë, *The Tenant of Wildfell Hall*

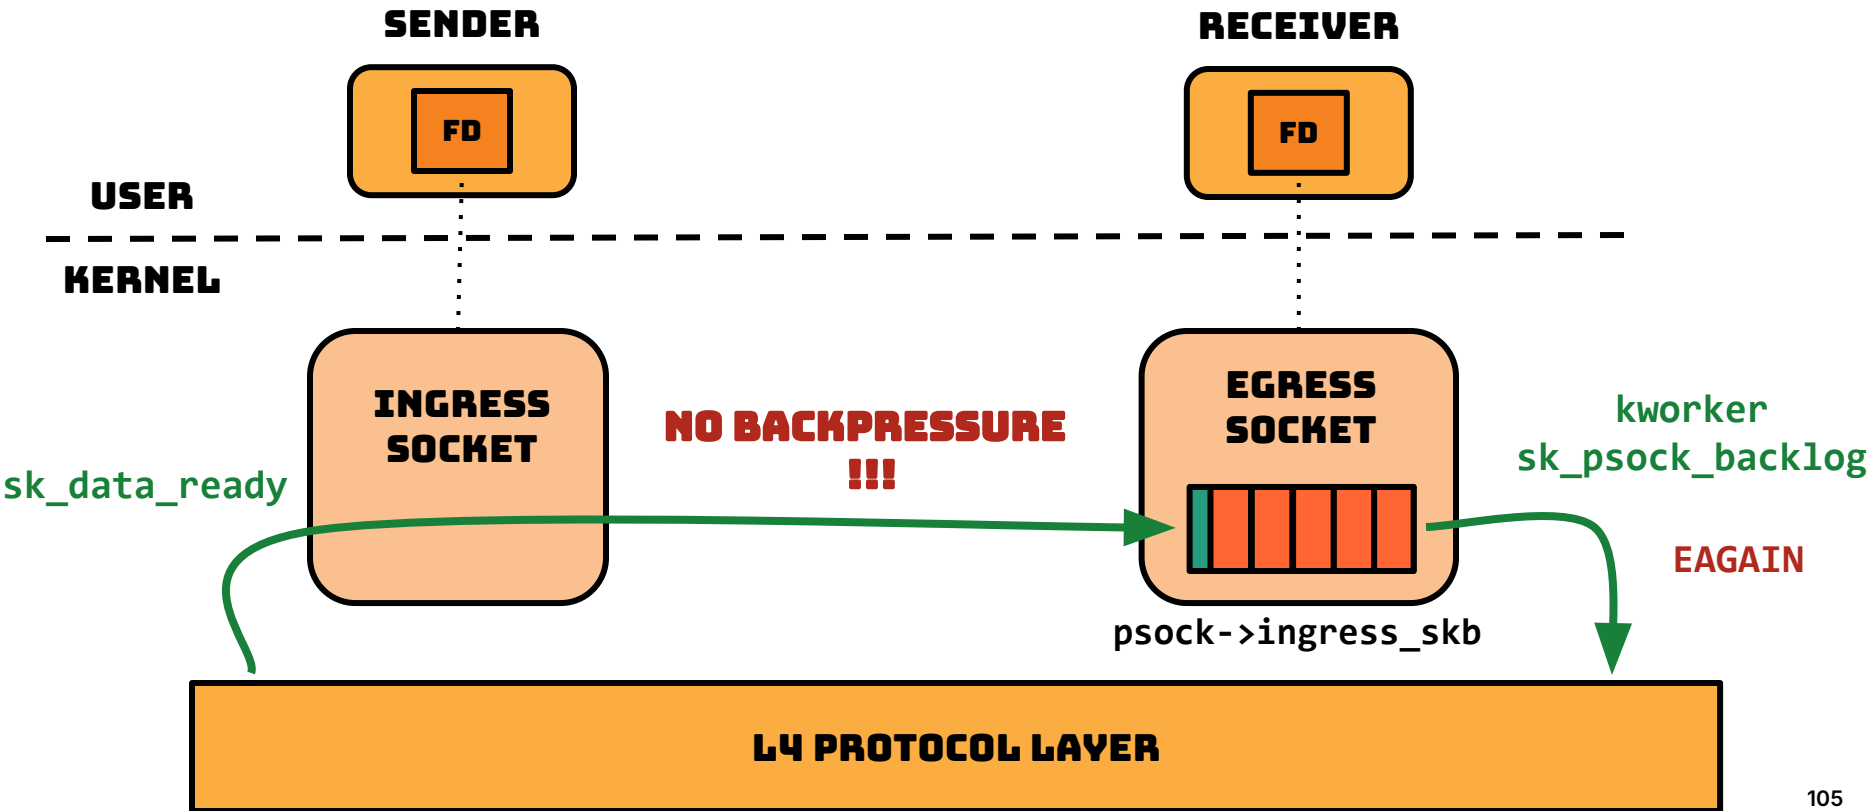


Redirect → ingress to egress → Internals





Redirect → ingress to egress → Internals



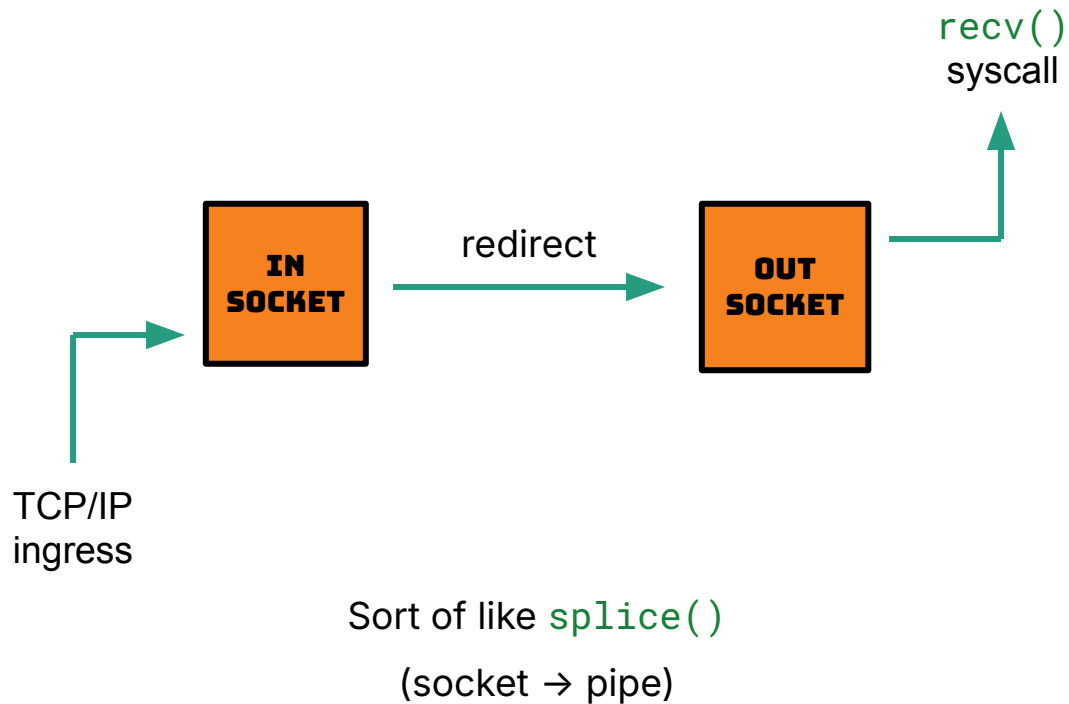
Redirect



ingress to local




Redirect → ingress to local





Redirect → ingress to local → How?



`BPF_PROG_TYPE_SK_SKB` prog

→ attached to `BPF_SK_SKB_VERDICT` hook

→ calls `bpf_sk_redirect_hash/map()` with `BPF_F_INGRESS` flag

→ returns `SK_PASS`







Redirect → ingress to local → What?

IN → OUT	TCP	UDP	UNIX STR	UNIX DGR	VSOCK STR	VSOCK SEQ
TCP	●	●	●	●	●	●
UDP	●	●	●	●	●	●
UNIX STR	●	●	●	●	●	●
UNIX DGR	●	●	●	●	●	●
VSOCK STR	●	●	●	●	●	●
VSOCK SEQ	●	●	●	●	●	●

any to any but VSOCK

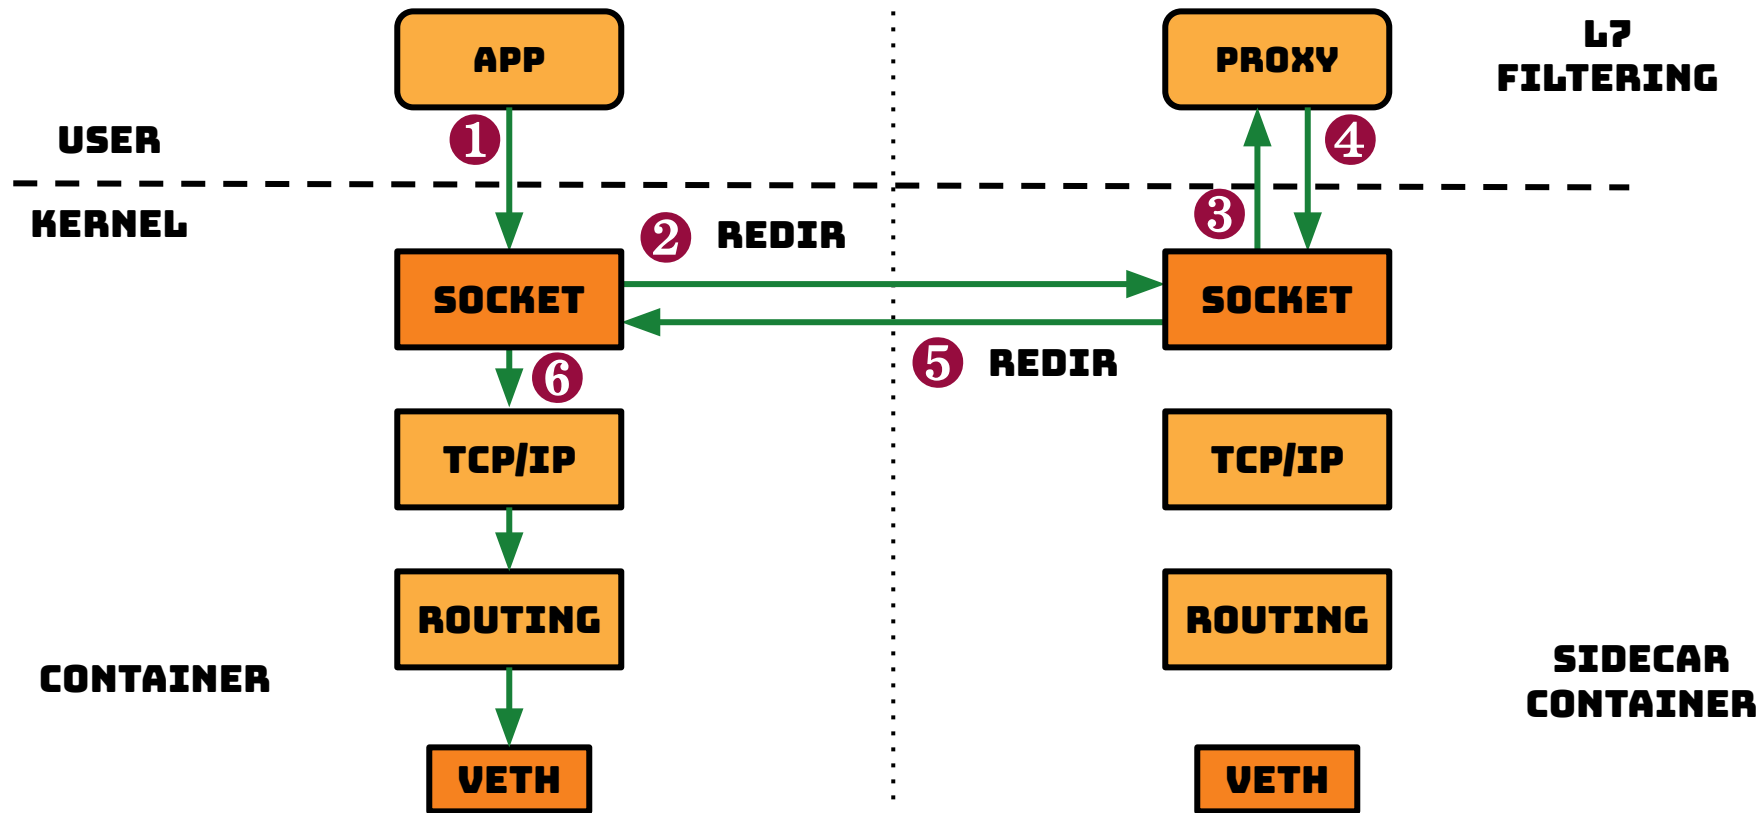
Cheatsheet - Redirect with SOCKMAP

redirect scenario	program type BPF_PROG_TYPE_*	attach type BPF_*	redirect helper function	redirect helper flags	in socket type	out socket type
send to local 	SK_MSG	SK_MSG_VERDICT	bpf_msg_redirect_*	BPF_F_INGRESS	TCP	any but VSOCK
send to egress 	SK_MSG	SK_MSG_VERDICT	bpf_msg_redirect_*	none	TCP	TCP
ingress to egress 	SK_SKB	SK_SKB_VERDICT	bpf_sk_redirect_*	none	any	any
ingress to local 	SK_SKB	SK_SKB_VERDICT	bpf_sk_redirect_*	BPF_F_INGRESS	any	any but VSOCK

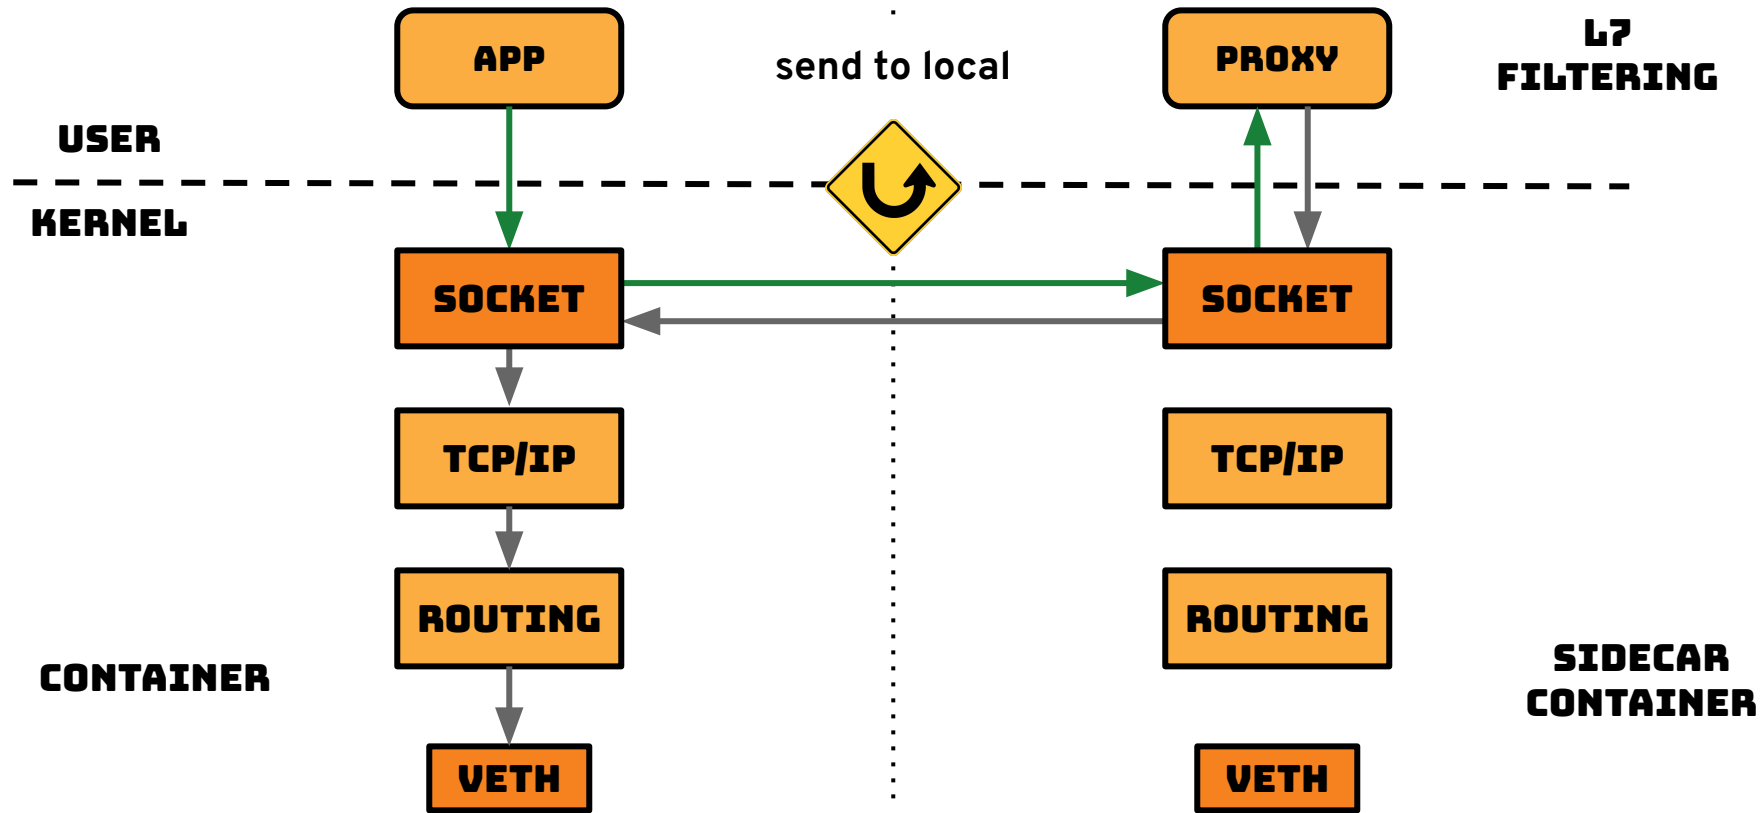
Real life use-cases



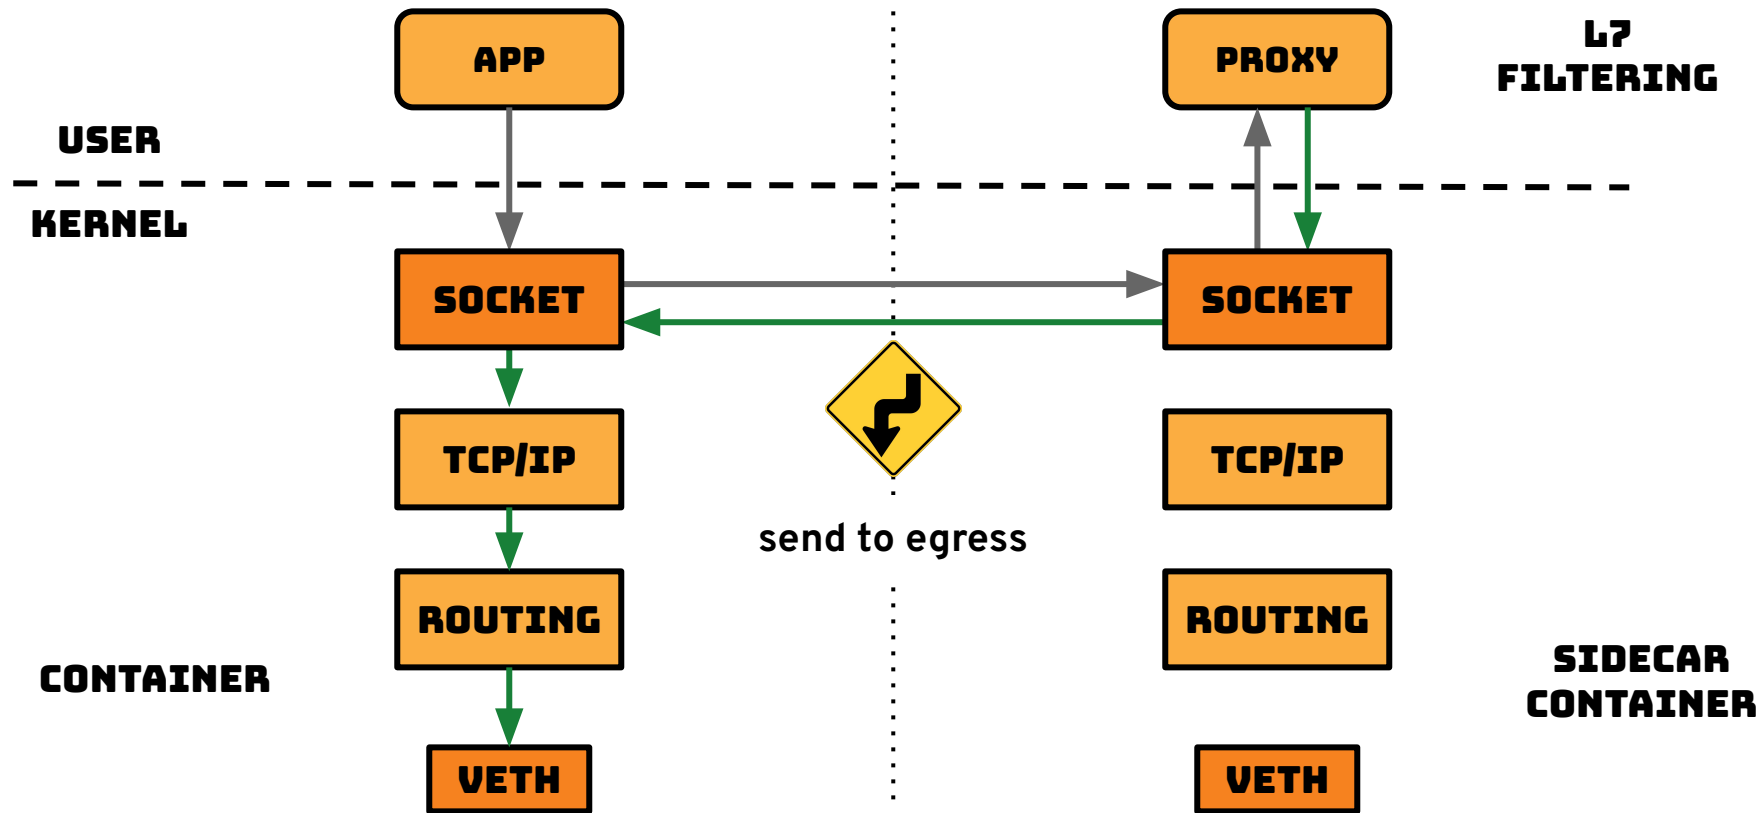
Cilium project (CNI for K8S)




Cilium project (CNI for K8S)



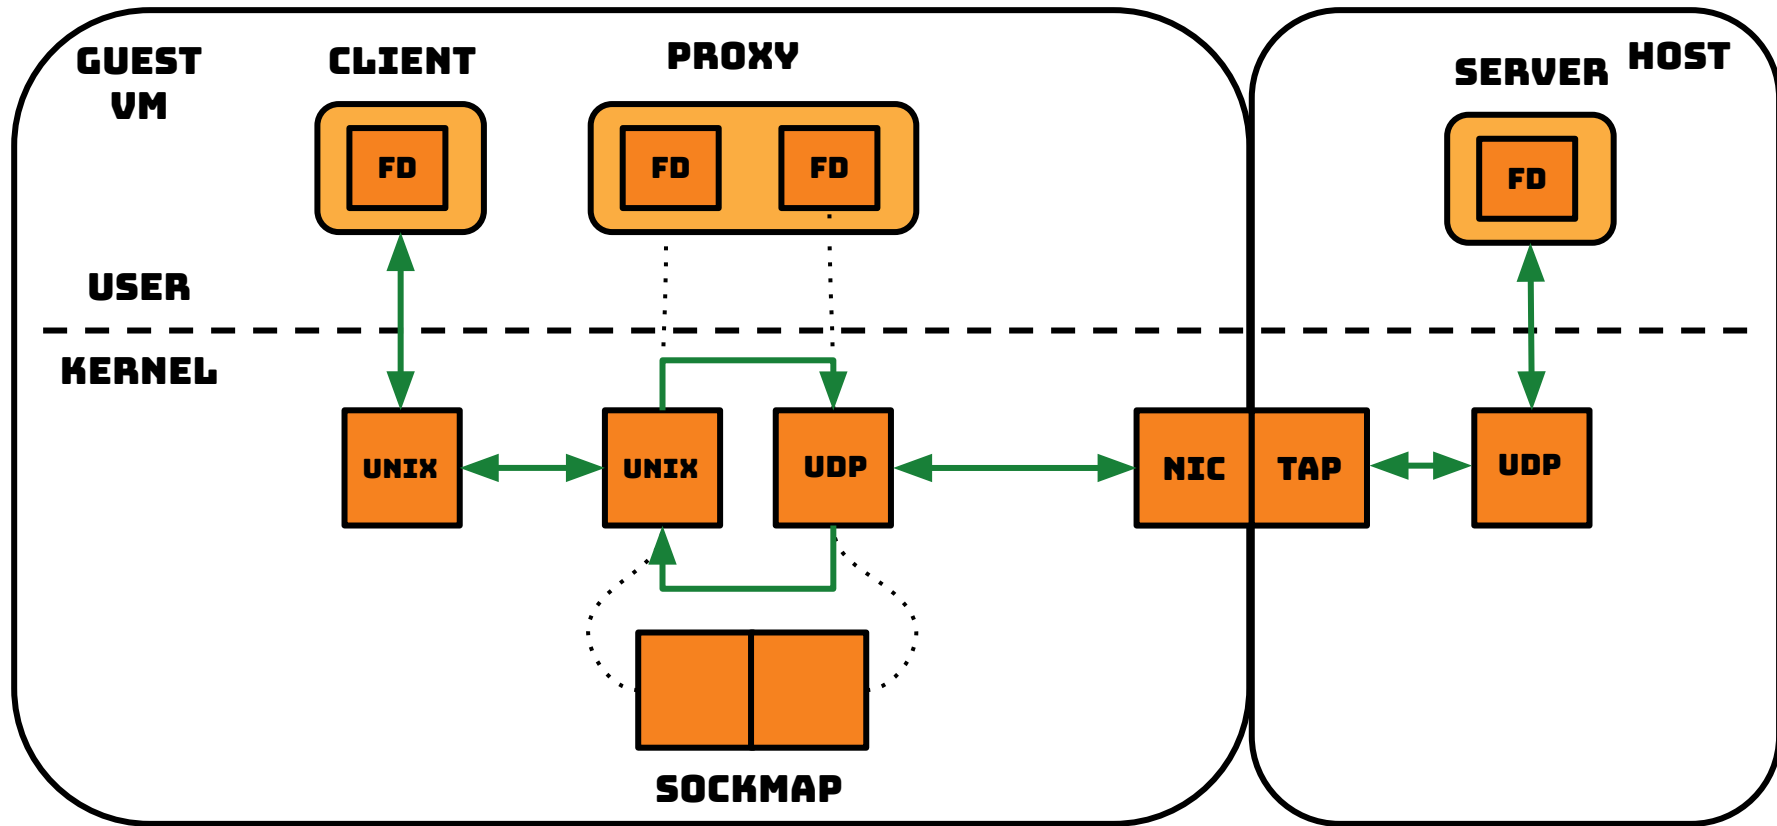
Cilium project (CNI for K8S)



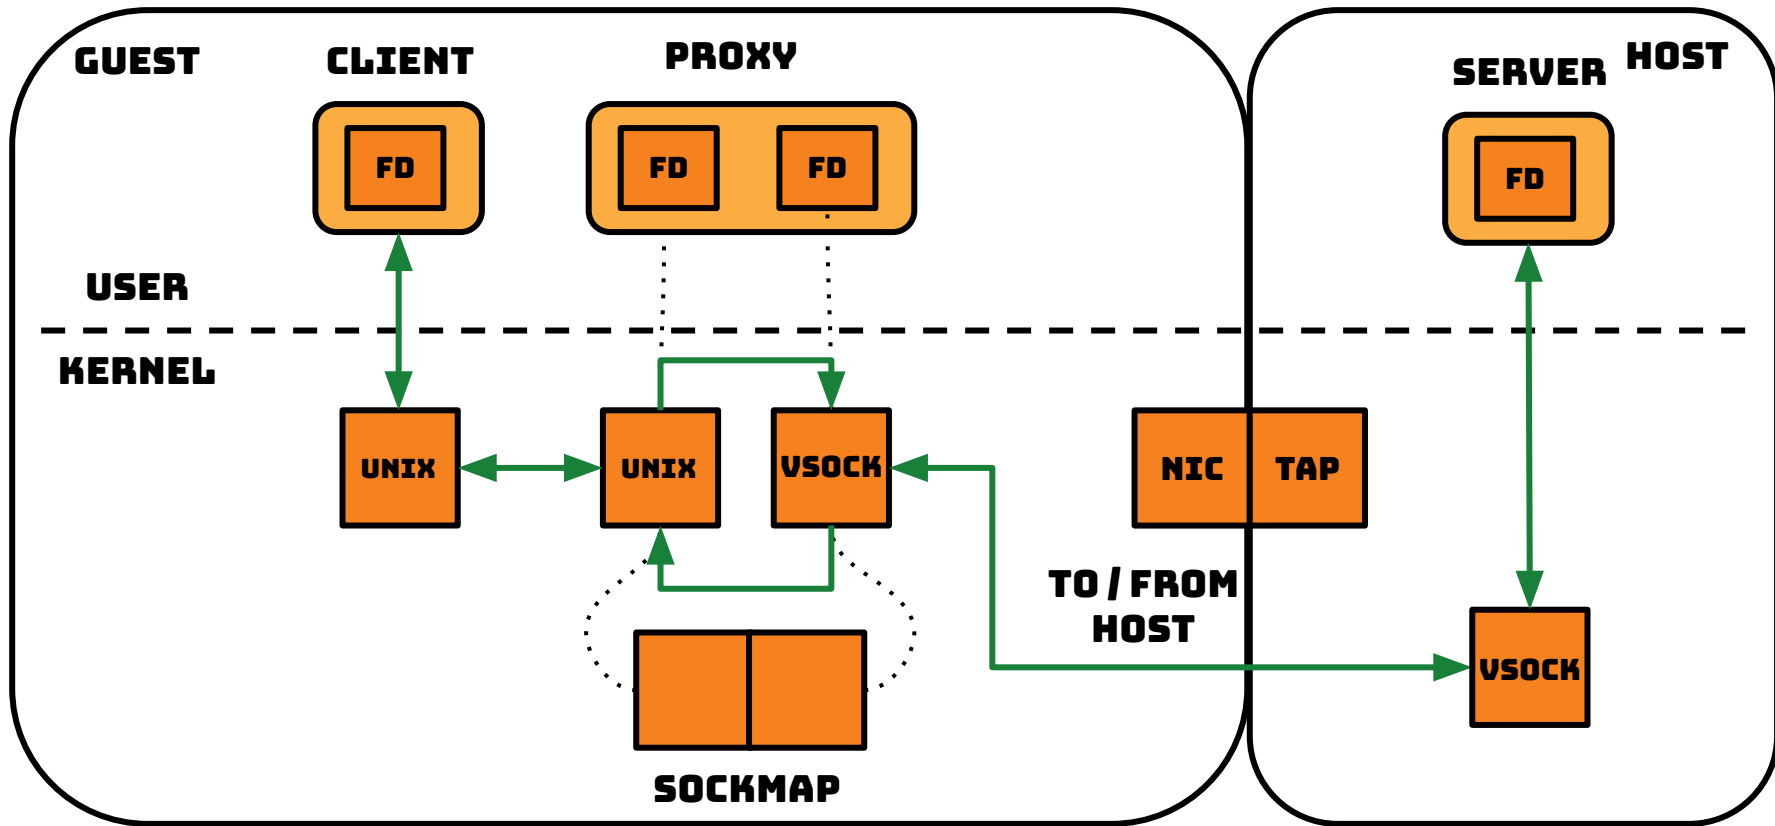
Bytedance (TikTok)



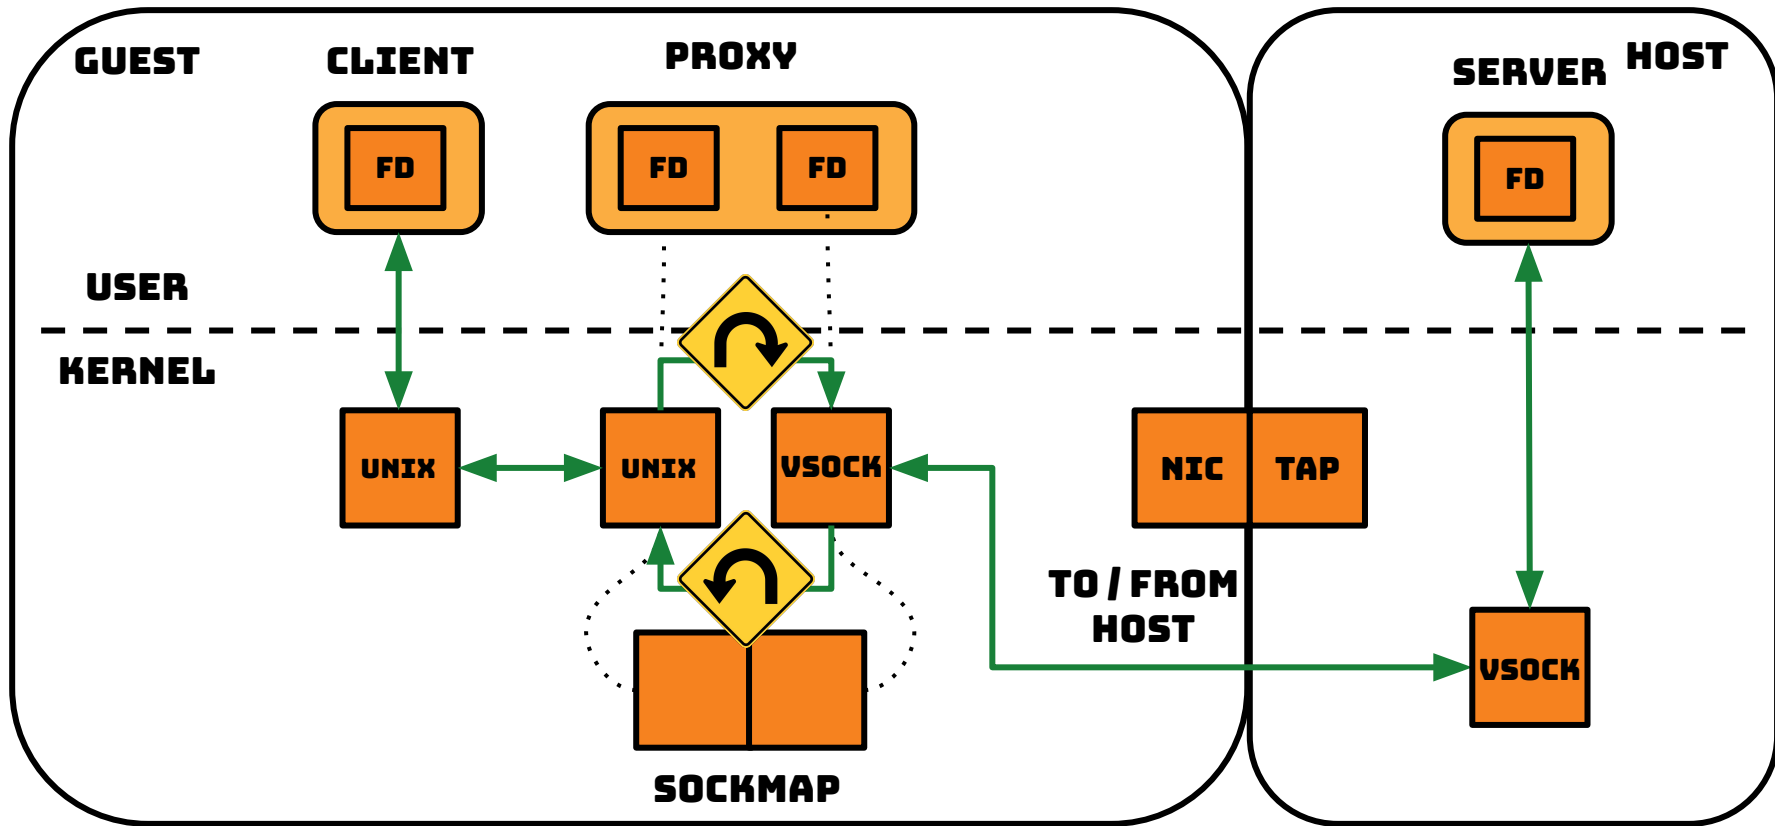
Bytedance (TikTok)



Bytedance (TikTok) → Improved



Bytedance (TikTok) → Improved v2



Where to learn more?





RESOURCES

1) **Linux Kernel → BPF Documentation → SOCKMAP and SOCKHASH map**

See unit tests with API usage examples

2) **LPC 2018: Combining kTLS and BPF for Introspection and Policy Enforcement**

See Daniel & John talk about Cilium SOCKMAP + kTLS use case ([video](#), [slides](#), [paper](#))

3) **Cloudflare Blog: SOCKMAP - TCP splicing of the future**

Read Marek review SOCKMAP from L7 proxy perspective

4) **eBPF Summit 2020: Steering connections to sockets with BPF socket lookup hook**

Another use case for SOCKMAP as a container ([video](#), [slides](#), [code](#))

THANK YOU



Reach out:

jakub@cloudflare.com

jsitnicki @ LinkedIn

Mailing lists:

bpf@vger.kernel.org

netdev@vger.kernel.org



code & slides repo



<https://github.com/jsitnicki/sockmap-project>

Attribution

- [Uncle Sam Wants You](#), Public Domain Pictures, CC0
- [A bunch of question marks](#), Wikimedia Commons, CC BY-SA 4.0 DEED
- [Nf knots](#), getarchive.net, CC0
- [Construction blocks](#), vectorportal.com, CC BY 4.0 DEED
- [Container crane image](#), vectorportal.com, CC BY 4.0 DEED
- [Golden Gate bridge silhouette](#), rawpixel.com, CC0
- [Evolution-des-wissens.jpg](#), Wikimedia Commons, CC BY-SA 3.0 DEED
- [Leather pouch](#), rawpixel.com, CC0
- [Redirection.svg](#), Wikimedia Commons, CC BY 3.0 DEED